

HYBRID-BERT AND METADATA DEEP LEARNING MODEL FOR TWITTER BOT DETECTION

*A Project Report Submitted in the Partial Fulfillment of
The Requirements for The Award of The Degree*

BACHELOR OF TECHNOLOGY IN COMPUTER SCIENCE AND ENGINEERING

Submitted By

SHAIK CHINNA MASTANVALI (22471A05P0)

Under the esteemed guidance of

SHAIK KHAJA MOHIDDIN BASHA, M. Tech

Assistant Professor



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

NARASARAOPETA ENGINEERING COLLEGE: NARASARAOPETA

(AUTONOMOUS)

**Accredited by NAAC with A+ Grade and NBA under Tier-1 and an ISO
9001:2015 Certified**

**Approved by AICTE, New Delhi, Permanently Affiliated to JNTUK, Kakinada
KOTAPPAKONDA ROAD, YALAMANDA VILLAGE, 522601**

2025-2026

NARASARAOPETA ENGINEERING COLLEGE
(AUTONOMOUS)
DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING



CERTIFICATE

This is to certify that the project that is entitled with the name **“HYBRIDBERT AND METADATA DEEP LEARNING MODEL FOR TWITTER BOT DETECTION”** is Bonafide work done by **SHAIK CHINNA MASTANVALI (22471A05P0)** in partial fulfillment of the requirements for the award of the degree of **BACHELOR OF TECHNOLOGY** in the Department of **COMPUTER SCIENCE AND ENGINEERING** during 2025-2026.

PROJECT GUIDE

Shaik Khaja Mohiddin Basha M.Tech
Assistant Professor

PROJECT CO-ORDINATOR

Syed Rizwana, B.Tech., M.Tech., (Ph. D).
Assistant Professor

HEAD OF THE DEPARTMENT

Dr.S. N. Tirumala Rao, M.Tech., Ph.D.
Professor & HOD

EXTERNAL EXAMINER

DECLARATION

I declare that this project work titled “**HYBRIDBERT AND METADATA USING DEEP LEARNING MODELS FOR TWITTER BOT DETECTION**” is composed by me that the work contains here is my own except where explicitly stated otherwise in the text and that this work had been submitted for any degree or professional qualification except as specified

SHAIK CHINNA MASTANVALI (22471A05P0)

ACKNOWLEDGEMENT

I wish to express my thanks to various personalities who are responsible for the completion of my project. I am extremely thankful to my beloved chairman, **Sri M. V. Koteswara Rao, B.Sc.**, who took keen interest in me in every effort throughout this course. I owe my sincere gratitude to my beloved principal, **Dr. S. Venkateswarlu, Ph.D.**, for showing his kind attention and valuable guidance throughout the course.

I express my deep-felt gratitude towards **Dr. S. N. Tirumala Rao, M.Tech., Ph.D.**, HOD of the CSE department, and also to my guide, **Shaik Khaja Mohiddin Basha, B.Tech., M.Tech. Assistant Professor** of the CSE department, whose valuable guidance and unstinting encouragement enabled me to accomplish my project successfully in time.

I extend my sincere thanks to **Syed Rizwana, B.Tech., M.Tech., (Ph. D)** Assistant Professor & Project Coordinator of the project, for extending his encouragement. Their profound knowledge and willingness have been a constant source of inspiration for me throughout this project work.

I extend my sincere thanks to all the other teaching and non-teaching staff in the department for their cooperation and encouragement during my B.Tech. degree. I have no words to acknowledge the warm affection, constant inspiration, and encouragement that I received from my parents. I affectionately acknowledge the encouragement received from my friends and those who were involved in giving valuable suggestions and clarifying my doubts, which really helped me in successfully completing my project.



INSTITUTE VISION AND MISSION

INSTITUTION VISION

To emerge as a Centre of excellence in technical education with a blend of effective student centric teaching learning practices as well as research for the transformation of lives and community.

INSTITUTION MISSION

M1: Provide the best class infra-structure to explore the field of engineering and research

M2: Build a passionate and a determined team of faculty with student centric teaching, imbining experiential, innovative skills

M3: Imbibe lifelong learning skills, entrepreneurial skills and ethical values in students for addressing societal problems



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

VISION OF THE DEPARTMENT

To become a center of excellence in nurturing the quality Computer Science & Engineering professionals embedded with software knowledge, aptitude for research and ethical values to cater to the needs of industry and society.

MISSION OF THE DEPARTMENT

The department of Computer Science and Engineering is committed to **M1:** Mould the students to become Software Professionals, Researchers and Entrepreneurs by providing advanced laboratories.

M2: Impart high quality professional training to get expertize in modern software tools and technologies to cater to the real time requirements of the Industry.

M3: Inculcate team work and lifelong learning among students with a sense of societal and ethical responsibilities.

Program Specific Outcomes (PSO's)

PSO1: Apply mathematical and scientific skills in numerous areas of Computer Science and Engineering to design and develop software-based systems.

PSO2: Acquaint module knowledge on emerging trends of the modern era in Computer Science and Engineering

PSO3: Promote novel applications that meet the needs of entrepreneur, environmental and social issues.



Program Educational Objectives (PEO's)

The graduates of the programmer are able to:

PEO1: Apply the knowledge of Mathematics, Science and Engineering fundamentals to identify and solve Computer Science and Engineering problems.

PEO2: Use various software tools and technologies to solve problems related to the academia, industry and society.

PEO3: Work with ethical and moral values in the multi-disciplinary teams and can communicate effectively among team members with continuous learning.

PEO4: Pursue higher studies and develop their career in software industry.

Program Outcomes:

PO1: Engineering knowledge: Apply the knowledge of mathematics, science, engineering fundamentals, and an engineering specialization to the solution of complex engineering problems.

PO2: Problem analysis: Identify, formulate, research literature, and analyze complex engineering problems reaching substantiated conclusions using first principles of mathematics, natural sciences, and engineering sciences.

PO3: Design/development of solutions: Design solutions for complex engineering problems and design system components or processes that meet the specified needs with appropriate consideration for the public health and safety, and the cultural, societal, and environmental considerations.

PO4: Conduct investigations of complex problems: Use research-based knowledge and research methods including design of experiments, analysis and interpretation of data, and synthesis of the information to provide valid conclusions.

PO5: Modern tool usage: Create, select, and apply appropriate techniques, resources, and modern engineering and IT tools including prediction and modeling to complex engineering activities with an understanding of the limitations.

PO6: The engineer and society: Apply reasoning informed by the contextual knowledge to assess societal, health, safety, legal and cultural issues and the consequent responsibilities relevant to the professional engineering practice.

PO7: Environment and sustainability: Understand the impact of the professional engineering solutions in societal and environmental contexts, and demonstrate the knowledge of, and need for sustainable development.

PO8: Ethics: Apply ethical principles and commit to professional ethics and responsibilities and norms of the engineering practice.

PO9: Individual and team work: Function effectively as an individual, and as a member or leader in diverse teams, and in multidisciplinary settings.

PO10: Communication: Communicate effectively on complex engineering activities with the engineering community and with society at large, such as, being able to comprehend and write effective reports and design documentation, make effective presentations, and give and receive clear instructions.

PO11: Project management and finance: Demonstrate knowledge and understanding of the engineering and management principles and apply these to one's own work, as a member and leader in a team, to manage projects and in multidisciplinary environments.

Project Course Outcomes (CO'S):

CO421.1: Analyze the System of Examinations and identify the problem.

CO421.2: Identify and classify the requirements.

CO421.3: Review the Related Literature

CO421.4: Design and Modularize the project

CO421.5: Construct, Integrate, Test and Implement the Project.

CO421.6: Prepare the project Documentation and present the Report using appropriate method.

Course Outcomes – Program Outcomes mapping

	PO1	PO2	PO3	PO4	PO5	PO6	PO7	PO8	PO9	PO10	PO11	PSO1	PSO2	PSO3
C421.1		✓										✓		
C421.2	✓		✓		✓							✓		
C421.3				✓		✓	✓	✓				✓		
C421.4			✓			✓	✓	✓				✓	✓	
C421.5					✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
C421.6									✓	✓	✓	✓		

Course Outcomes – Program Outcome correlation

	PO1	PO2	PO3	PO4	PO5	PO6	PO7	PO8	PO9	PO10	PO11	PSO1	PSO2	PSO3
C421.1	2	3										2		
C421.2			2		3							2		
C421.3				2		2	3	3				2		
C421.4			2			1	1	2				3	2	
C421.5					3	3	3	2	3	2	2	3	2	1

Note: The values in the above table represent the level of correlation between CO's and POs:

1. Low level
2. Medium level
3. High level

Project mapping with various courses of Curriculum Attained POs:

Name of the course from which principles are applied in this project	Description of the device	Attained PO
C2204.2, C22L3.2	Gathering the requirements and defining the problem, plan to develop model for detection and classification of OSCC	PO1, PO3, PO8
CC421.1, C2204.3, C22L3.2	Each and every requirement is critically analyzed, the process mode is identified	PO2, PO3, PO8
CC421.2, C2204.2, C22L3.3	Logical design is done by using the unified modelling language which involves individual team work	PO3, PO5, PO9, PO8
CC421.3, C2204.3, C22L3.2	Each and every module is tested, integrated, and evaluated in our project	PO1, PO5, PO8
CC421.4, C2204.4, C22L3.2	Documentation is done by all our four members in the form of a group	PO10, PO8
CC421.5, C2204.2, C22L3.3	Each and every phase of the work in group is presented periodically	PO8, PO10, PO11
C2202.2, C2203.3, C1206.3, C3204.3, C4110.2	Implementation is done and the project will be handled by the social media users and in future updates in our project can be done based on detection for Oral Cancer	PO4, PO7, PO8

ABSTRACT

This paper presents a deep learning method to detect Twitter bots by combining two types of user information: tweet content and behavioral data. The proposed model uses BERT, a transformer-based language model, to capture patterns in user tweets. It also includes a lightweight neural network that processes structured details like follower count, posting frequency, and verification status. By merging textual content with user behavior features, the model effectively tells apart human-operated accounts from automated (bot) accounts. The system trains and tests on the TwiBot-20 dataset, which includes a wide range of annotated Twitter profiles. To address class imbalance during training, we use Focal Loss. We optimize the model using the Adam optimizer with a low learning rate over five epochs. Our method shows impressive results, achieving over 94% accuracy and an F1- score of 0.935. We verify the model's strength through thorough evaluation using ROC curves and confusion matrices. Additionally, the system allows for real-time prediction, making it ideal for large-scale content moderation. The model's design is also open to future improvements, such as multilingual bot detection.

INDEX

SNO	CONTENT	PAGENO
1	INTRODUCTION 1.1 MOTIVATION 1.2 PROBLEM STATEMENT 1.3 OBJECTIVES	1 4 5 6
2	LITERATURE SURVEY	7
3	SYSTEM ANALYSIS 3.1 EXISTING SYSTEM 3.1.1 DISADVANTGES OF DETCTION 3.2 PROPOSED SYSTEM 3.3 FEASIBILITY STUDY	9 9 11 12 14
4	SYSTEM REQRIMENTS 4.1 SOFTWARE REQRIMENTS 4.2 REQRIMENT ANALYSIS 4.3 HARDWARE REQRIMENTS 4.4 SOFTWARE 4.5 SOFTWARE DESCRIPTION	16 16 17 18 18 19
5	SYSTEM DESIGN 5.1 SYSTEM ARCHITECUTRE 5.2 DATASET 5.3 DATA PREPROCESSING 5.4 FEATURE EXTRACTION 5.5 MODEL BUILDING 5.6 CLASSIFICATION 5.7 UML DIAGRAM	20 20 21 23 25 27 29 31
6	IMPLEMENTATION	33
7	RESULT AND ANALYSIS	62
8	TEST CASES	65
9	USER INTERFACES	66
10	CONCLUSIONS	68
11	FUTURE SCOPE	69
12	REFERENCES	70
13	RESEARCH ARTICLE	73

LIST OF FIGURES

Sno	List Of Figures	Pageno
1	FIG:1 REPRESENTATION OF HYBIRDBERT+METADATA FUSION	2
2	FIG. 3.1. FLOW CHART OF THE TWITTER BOT DETECTION	10
3	FIG 3.2. FLOW CHART OF PROPOSED SYSTEM	12
4	FIG 5.2 DATASET ARCHITECTURE OF PROPOSED MODEL	22
5	FIG 5.6.5 CLASSIFICATION OVERVIEW OF PROPOSED MODEL	30
6	FIG 5.7 OVERVIEW OF UML DIAGRAM OF PROPOSED MODEL	31
7	FIG 7.A MODEL COMPARSION OF DIFFERENT MODELS	62
8	FIG 7.1 CONFUSION MATRIX OF HUMAN AND BOTS	63
9	FIG 7.2 TRAINING PROGRESS OF LOSS AND EPOCH	63
10	FIG 7.3 T-SNE VISUALIZATION OF MODEL	64
11	FIG 8.1.1 HUMAN TWEET OUTPUT	65
12	FIG 8.1.2 BOT TWEET OUTPUT	65
13	FIG 9.1 USERS HOME SCREEN	66
14	FIG 9.2 USERS PREDICTION SCREEN	66
15	FIG 9.3 USERS FEEDBACK SCREEN	67
16	FIG 9.4 USERS CONTACT FORM SCREEN	67

1. INTRODUCTION

Social media platforms such as Twitter have increasingly become targets for automated accounts called bots. These accounts pose a serious threat to the credibility and reliability of digital communication. Their diverse behaviors, adaptive strategies, and varying degrees of influence make them challenging to detect. While some bots perform routine or harmless tasks like posting weather updates or sharing news, others are designed to mislead users, spread false information, and manipulate online discussions. Therefore, distinguishing bots from real users is essential to preserve the integrity of online spaces, maintain information credibility, and foster trust across digital communities. Traditional detection techniques that depend solely on text-based features often struggle with accuracy, as modern bots can generate human-like content. This highlights the need for more advanced hybrid approaches that integrate both textual and behavioral metadata-based cues to enhance detection performance.

The growing influence of bots operates on a global scale and affects individuals, organizations, and even government systems. Although some automated accounts contribute positively to information dissemination, malicious ones exploit automation to amplify spam, misinformation, and propaganda. The rapid growth of these entities can be attributed to easily available automation tools, improved account-generation technologies, and insufficient defense mechanisms. Public awareness regarding these bots remains limited, resulting in delayed detection and prolonged influence.

Studies indicate that automated accounts make up an estimated 9 to 15 percent of Twitter's total users. Their impact becomes particularly concerning during sensitive events such as elections, global crises, or public movements when misinformation can significantly alter public opinion and decision-making. Among these, political bots are considered the most disruptive, as they operate in coordination, post frequently, and distort online narratives at scale.

A critical issue in combating bots lies in the absence of proactive detection. Users often struggle to differentiate authentic accounts from automated ones, and platforms typically intervene only after harmful activity is widespread.

patterns in user metadata. Swift and accurate identification of these signals is essential to contain the potential harm caused by bots. Another major challenge arises from the growing sophistication of malicious bots. They now imitate human behavior convincingly, engaging in trending topics, interacting with real users, and producing contextually relevant posts to avoid detection. While large organizations and research institutions are developing complex detection frameworks, smaller entities and individual users often lack access to such tools. Furthermore, the absence of standardized detection protocols and slow platform responses give harmful bots more time to manipulate discussions and spread misinformation.

To address these limitations, hybrid detection frameworks have emerged as a promising direction. By combining text-based semantic analysis with structured metadata, these models leverage the strengths of both natural language understanding and machine learning classification. The present research introduces a Hybrid BERT + Metadata model that integrates contextual embeddings from BERT with behavioral insights from user metadata through a Multi-Layer Perceptron network. The approach employs advanced preprocessing techniques, including text tokenization, metadata normalization, and class balancing, to ensure efficient data representation and improved classification accuracy.

The proposed hybrid framework aims to deliver higher precision and reliability in distinguishing between bots and human users compared to conventional single-input systems. While BERT captures linguistic nuances from tweet content, the metadata branch enhances discrimination




FIG:1 REPRESENTATION OF HYBIRDBERT+METADATA FUSION

The **Hybrid BERT + Metadata Model** is designed to leverage both textual semantics and user behavior attributes to enhance the accuracy of Twitter bot detection. It operates through three key processing stages:

1. Textual Feature Extraction (Left Module)

In this stage, the textual data—such as user tweets, profile descriptions, or recent activity—is first **tokenized** into input IDs, attention masks, and token type identifiers. These structured inputs are then passed through the **Bidirectional Encoder Representations from Transformers (BERT)** model, which captures deep contextual and semantic relationships between words. The output from BERT is a **768-dimensional embedding vector**, specifically derived from the [CLS] token, which represents the overall semantic meaning of the text segment.

2. Metadata Feature Extraction (Right Module)

The metadata stream focuses on user account-level attributes such as followers count, these behavioral features are fed into a Multi-Layer Perceptron (MLP) network, which prevent overfitting. The MLP transforms these numerical inputs into a 896-dimensional feature vector, effectively capturing the behavioral patterns associated user activity.

3. Fusion and Final Classification

After processing both streams, the **BERT embeddings** and the **MLP metadata features** are concatenated to form a unified representation. This fused vector integrates both linguistic context and behavioural evidence, which is then passed into a **fully connected classification layer**. The final output of this model is a **binary prediction**, identifying whether the Twitter account is operated by a **human** or a **bot**.

Once both the textual and behavioural features are obtained, they are merged into a single hybrid feature space. This fusion enables the model to assess patterns that neither textual nor metadata features alone could reveal. The combined features are then passed to a classification layer, which learns to differentiate between genuine human users and automated bot profiles. Through this dual-stream architecture, the model achieves a balance between semantic comprehension and behavioural analysis, resulting in higher precision and robustness in detecting bots operating under varied strategies or disguises.

1.1 MOTIVATION

Twitter bots have emerged as a significant challenge in maintaining the authenticity, security, and reliability of online communication. These automated accounts disrupt information exchange, manipulate discussions, and erode public trust. Detecting such accounts promptly and accurately is crucial to ensuring safe digital spaces. Traditional manual methods such as observing user activity, posting frequency, or follower patterns are slow, inconsistent, and largely ineffective, especially as bots increasingly mimic genuine human interactions.

Given the vast and evolving nature of social media data, automated detection systems are necessary to efficiently analyze both textual and behavioral patterns. Conventional models that rely only on text analysis or metadata fail to capture the full complexity of online behavior, leading to limited detection accuracy. This limitation highlights the need for hybrid models that can integrate multiple sources of information to improve precision and robustness.

The objective of this study is to develop an intelligent and scalable detection framework that minimizes manual intervention, reduces misclassification, and enhances the reliability of identifying automated accounts. The proposed hybrid model combines the strengths of two complementary components BERT which extracts deep contextual features from tweets and profile descriptions and a Multi-Layer Perceptron MLP that interprets structured metadata such as follower count, friend count, listed count, status updates, and verification status. While BERT captures semantic relationships and linguistic patterns, the MLP component focuses on behavioral indicators, and their fusion creates a comprehensive decision-making mechanism for classifying accounts as human or bot.

To enhance practical usability, the system can be implemented as an interactive tool where users, researchers, or organizations can input account details to receive real time classification feedback. This setup supports quicker identification of harmful automated profiles, assists in managing misinformation, and contributes to a more trustworthy digital environment. The proposed framework ultimately delivers a reliable, scalable, and effective solution for bot detection by unifying advanced text representation and metadata driven behavioral analysis.

1.2 PROBLEM STATEMENT

Twitter bots have become a growing problem that affects the credibility, safety, and openness of social platforms. These automated accounts can change how information spreads and how people interact online. Their influence depends on how often they post, how many followers they have, and the goals of those who create them. Harmful bots can shape opinions, spread false information, and interrupt real conversations, causing effects that go beyond social media.

The scale of a bot's impact is linked to how active and well-connected it is, as well as how realistic its behavior appears. Some may seem harmless at first but later take part in organized efforts to mislead people or push certain agendas. Such actions can change the direction of political debates, health awareness, or social campaigns. Detecting these bots early is very important because once their content spreads widely, it becomes hard to control or undo the damage.

Finding and classifying bots is challenging because they come in many forms. Some post very frequently, while others stay hidden and act carefully within specific groups. Many use advanced techniques to look like human users by copying normal writing styles, engagement habits, and timing patterns.

If left unchecked, bots can cause serious harm by spreading rumors, creating confusion, and influencing major events such as elections or public discussions. Their overall effect depends on how large their networks are and how long they remain active. Coordinated groups of political bots, for example, can flood conversations and silence real users, giving a false sense of public opinion.

To handle this problem, detection systems must be both quick and dependable. Because bots continue to evolve, simple rule-based methods no longer work. A better approach is to analyze both text and behavior together. By combining language features from tweet content with numerical details such as followers, post counts, and account verification, hybrid detection models can improve accuracy. This method helps identify harmful accounts sooner, reduce misinformation, and keep social media spaces more genuine and trustworthy.

1.3 OBJECTIVE

The Twitter Bot Detection and Classification project focuses on developing an intelligent system capable of accurately separating real users from automated accounts on Twitter. The framework combines two analytical layers: BERT, which interprets the language used in tweets, and a lightweight neural network that studies user metadata such as follower count, activity rate, and verification status. This dual approach allows the model to understand both textual meaning and behavioral tendencies, improving reliability and overall detection accuracy.

The system is deployed through an interactive web platform built with Flask, enabling users to analyze Twitter accounts or upload datasets for instant classification. Input validation ensures that only correct and usable information is processed, while users receive clear feedback when errors occur. The tool is intended to assist researchers, analysts, and organizations by simplifying detection tasks, reducing manual checks, and improving decision-making speed.

Scalability is built into the design so the model can later expand to identify different categories of bots, connect with live Twitter data streams, and continuously update its learning process. By emphasizing early and precise identification, the project enhances digital trust, limits the spread of misinformation, and supports a safer online communication space.

In essence, this work demonstrates how combining natural language processing with behavioral analysis can produce a practical and adaptable defense system against automated social media manipulation.

2. LITERATURE SURVEY

Twitter bot detection and classification using deep learning have attracted significant research attention in recent years, primarily due to the growing demand for reliable and automated systems to counter malicious activities on social platforms. Earlier attempts relied on traditional machine learning approaches such as Support Vector Machines and Random Forest classifiers to analyze user profiles and content features. While these techniques provided some success, they often struggled with challenges related to feature engineering, generalization, and adaptability.

In contrast, deep learning methods particularly transformer-based architectures like BERT and Roberta have shown remarkable improvements by learning contextual and semantic patterns directly from data. The rise of transfer learning has further advanced this domain, allowing large-scale pre-trained models to be adapted for social media datasets with limited annotations. Researchers have also explored multimodal strategies that combine textual information with metadata and interaction features, achieving improved detection performance compared to single-source models.

Yang and collaborators highlighted the advantages of transfer learning by employing transformer-based encoders for identifying bots, reporting significant improvements over traditional classifiers. Similarly, Imran proposed a hybrid design that integrates BERT text embeddings with structured metadata such as account creation date and follower statistics. This dual-stream approach demonstrated notable gains in accuracy compared to text-only baselines, emphasizing the value of multimodal integration.

Ahmad and Choudhury conducted a systematic review of machine learning and deep learning applications for bot detection, showing that pre-trained models such as BERT and Roberta consistently deliver performance levels above conventional approaches. However, they also pointed out challenges in fine-tuning strategies and computational requirements, which remain obstacles for widespread deployment.

Shah introduced a multimodal framework that fused metadata with tweet embeddings, resulting in enhanced robustness and improved classification across varied datasets. Lightweight architectures, combined with data augmentation and domain-specific preprocessing, have also been reported to improve scalability and real-time applicability.

Agrawal explored the use of attention-based BERT models for bot detection, while Sireesha demonstrated the importance of combining user profiles with text data for high-precision classification. These studies reinforce the significance of fusing behavioral and linguistic signals to capture sophisticated patterns of automated accounts.

Graph-based models have also emerged as a promising direction.

By leveraging retweet and follower networks, researchers have achieved higher reliability in distinguishing bots from humans. Radhi developed an efficient framework that incorporated activity-based features with network signals, reducing both false positives and detection time. Similarly, Tomaka and colleagues proposed a transformer-enhanced hybrid model, reporting robustness and improved generalization across multiple benchmarks.

Other works, such as those by Sharma and Gujral, revisited classical machine learning models, showing that when carefully combined with feature selection techniques, traditional classifiers can still deliver competitive results. Moturi and collaborators, while focusing on healthcare data, proposed hybrid feature extraction and ensemble learning techniques that can be adapted for bot detection to enhance decision-making accuracy.

Islam introduced unsupervised learning approaches using clustering and dimensionality reduction for improving detection efficiency. Amin and Sharif developed metadata-centered preprocessing pipelines that improved model interpretability and classification accuracy. Likewise, Abiwinanda designed deep neural networks with dropout and normalization techniques that enhanced model generalization. Setha and Raja also employed transformer-based models, demonstrating the relevance of transfer learning in domains with limited annotated datasets.

Hassan Ali Khan presented a lightweight hybrid model that outperformed large-scale pre-trained encoders while consuming fewer computational resources. Overall, Deep learning's scalability and adaptability make it a cornerstone for modern bot detection. Transformer-based encoders, when coupled with metadata and network features, have proven to be particularly effective. In addition, supporting methods such as transfer learning, domain-specific preprocessing, and data augmentation ensure reliable performance even in resource-constrained settings. This positions deep learning as one of the most powerful tools for combating social media automation and misinformation scale.

3. SYSTEM ANALYSIS

3.1 EXISTING SYSTEM

Twitter bot detection has traditionally relied on manual inspection of user profiles, tweet patterns, and interaction behaviors by moderators and domain experts. While human assessment can be effective in identifying suspicious activity, it is slow, inconsistent, and heavily dependent on expertise. Moreover, manual interpretation is prone to oversight, often leading to undetected bots that mimic human-like behavior. To address these challenges, computational approaches have been introduced, ranging from classical Machine Learning models to advanced Deep Learning architectures.

Early Machine Learning (ML) techniques such as Support Vector Machines (SVMs), Random Forests (RF), Logistic Regression, and k-Nearest Neighbors (k-NN) focused on using handcrafted features extracted from Twitter data. These features included follower- friend ratios, posting frequency, retweet patterns, and profile-based attributes. Although these models achieved reasonable accuracy (around 70–85%), their performance was constrained by feature engineering dependency, poor adaptability across datasets, and difficulty in capturing the complex behavioral traits of sophisticated bots.

The emergence of Deep Learning transformed bot detection by enabling models to learn latent representations directly from raw textual content. Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs) were initially applied to tweet classification tasks, achieving improved accuracy levels of 85–92%. However, these models faced challenges such as high computational demands, difficulty in handling long-range dependencies in language, and overfitting on limited datasets.

To overcome these limitations, transformer-based architectures such as BERT and Roberta have been widely adopted. Pretrained models fine-tuned on social media datasets demonstrated significant improvements by capturing semantic and contextual information in tweets. Transfer learning not only reduced training time but also enhanced accuracy by an additional 5–10% compared to conventional deep learning models. In addition to text-only models.




FIG. 3.1. FLOW CHART OF THE TWITTER BOT DETECTION

This flowchart (Fig. 3.1) illustrates a conventional pipeline for Twitter bot detection using text-based and machine learning approaches. The process starts with an input dataset containing user tweets and profile details. In the pre-processing stage, text is cleaned by removing stop words, special characters, and links, followed by tokenization and normalization. For user metadata, features such as follower count, following count, account age, and tweet frequency are extracted and standardized.

After pre-processing, feature extraction is performed. Traditional models rely on handcrafted features such as linguistic patterns, posting behavior, and network statistics. Feature selection techniques, including Principal Component Analysis (PCA) or statistical correlation methods, are used to reduce dimensionality and improve computational efficiency.

Finally, classifiers such as Support Vector Machines (SVM), Random Forests (RF), k- Nearest Neighbors (k-NN), or standalone Convolutional Neural Networks (CNNs) are applied to distinguish between human-operated and automated (bot) accounts. While these approaches provide reasonable performance.

3.1.1 DISADVANTAGES OF TWITTER BOT DETECTION

Despite significant advancements in automated Twitter bot detection, the existing systems face several limitations:

- **Dependence on Text-Only Features:**

Most models rely solely on tweet content, ignoring valuable metadata such as account age, follower/following ratio, and posting patterns. This limits their ability to detect sophisticated bots that mimic human-like text.

- **Requirement of Large Labeled Datasets:**

Deep learning models demand large-scale annotated datasets for effective training. However, creating labeled datasets of bot and human accounts is challenging, as annotation requires manual verification and is prone to error.

- **Overfitting on Small Datasets:**

Due to limited availability of benchmark datasets like TwiBot-20, models often overfit—performing well on training data but failing to generalize effectively to unseen accounts.

- **High Computational Complexity:**

Transformer-based architectures such as BERT and RoBERTa require significant GPU resources and memory, making them less practical for real-time deployment in resource-constrained environments.

- **Lack of Robustness Against Evolving Bot Strategies:**

Bots continuously adapt their behavior and language to evade detection. Existing text-only systems struggle to remain effective against new bot strategies, reducing long-term reliability.

- **Sensitivity to Noisy or Incomplete Data:**

User profiles and tweets often contain incomplete, noisy, or misleading information (e.g., fake bios, repetitive posts). Text-only classifiers are easily misled under such conditions.

- **Difficulty in Handling Class Imbalance:**

Datasets are often imbalanced, with fewer labeled bot accounts compared to human accounts. This causes models to be biased towards human classification, reducing bot detection accuracy.

3.2 PROPOSED SYSTEM

The proposed model employs a **hybrid framework** that combines **Bidirectional Encoder Representations from Transformers (BERT)** with a **metadata-driven Multi-Layer Perceptron (MLP)** to improve the accuracy and reliability of Twitter bot detection. BERT is applied to capture contextual and semantic patterns from textual content, while the metadata component incorporates profile-level attributes such as followers count, friends count, activity statistics, and verification status. By integrating these complementary sources of information, the model achieves stronger generalization and consistency in distinguishing between human-operated and automated accounts.




FIG 3.2. FLOW CHART OF PROPOSED SYSTEM

The workflow begins with **data preprocessing**, where user profile information and tweets are prepared for analysis. Textual data is cleaned by removing URLs, mentions, special characters, and stop words, ensuring that the input is consistent and meaningful. Metadata attributes, including followers count, friends count, statuses count, listed count, and verification status, are normalized to a uniform scale to prevent bias during training.

After preprocessing, the textual component is passed through **BERT**, which generates rich contextual embeddings that capture semantic and linguistic patterns from the tweets. In parallel, the **metadata features** are processed through a multi-layer

Perceptron (MLP), which learns non-linear relationships from structured user profile information.

The outputs from both branches are then **fused into a joint representation**, combining textual semantics with behavioral metadata. This integrated feature space is fed into a final classification layer, which determines whether a given account is human- operated or automated (bot). The hybrid design ensures a balanced contribution from both text- based and profile-based evidence.

Advantages over Existing Systems

- 1. Improved Accuracy and Reliability:** The integration of BERT for text analysis with metadata-based learning enhances the model's decision-making compared to using either component alone.
- 2. Comprehensive Feature Utilization:** By leveraging both tweet content and account metadata, the model captures multiple dimensions of user behavior, making detection more robust.
- 3. Better Generalization:** Normalization and balanced fusion reduce overfitting and allow the model to adapt effectively across different user types and datasets.
- 4. Context-Aware Detection:** Unlike models that rely only on surface-level patterns, BERT embeddings capture deep semantic meaning, improving classification of sophisticated bot accounts.
- 5. High Performance Metrics:** Strong F1-score, AUC, and precision-recall balance ensure reduced false positives (mislabeling humans as bots) and false negatives (failing to detect bots).
- 6. Scalable and Practical:** The framework can be extended to other social platforms with minimal modifications, providing a scalable solution for large-scale bot detection.

3.3 FEASIBILITY STUDY

1. Technical Feasibility

- **Automated Feature Extraction:**

BERT provides powerful contextual embeddings that capture semantic nuances from tweets, eliminating the need for manual feature engineering. These embeddings represent linguistic patterns and user intent more effectively, improving classification outcomes.

- **Complementary Metadata Utilization:**

While BERT handles unstructured textual data, the metadata module (MLP) processes structured account-level attributes such as followers, friends, activity count, and verification status. The combination strengthens the model by integrating both behavioral and linguistic perspectives.

- **Improved Accuracy and Generalization:**

The hybrid design mitigates overfitting risks by balancing textual and metadata inputs. Metadata contributes robustness where text signals are weak, while BERT enhances semantic understanding, resulting in better overall accuracy and generalization.

- **Scalability:**

The model is adaptable to various social media platforms beyond Twitter, as long as user text and account-level features are available. It can scale efficiently to large datasets through modern GPU-based training pipelines.

2. Operational Feasibility

- **Interpretability:** Although transformer-based models like BERT are often viewed as complex, the addition of a metadata-based MLP provides more transparency by highlighting which account-level features (e.g., follower-following ratio, activity frequency, verification status) contribute to.

- **Data Handling:**

The model requires labeled Twitter accounts (bot/human) for training but is flexible enough to handle varying text lengths and heterogeneous metadata. It can process multi-modal inputs—tweets and profile features—making classification more reliable across different user types.

- **Maintenance and Upgradation:**

The framework can be updated seamlessly with new training data to reflect evolving bot behaviors. Retraining only the classification head or metadata component is sufficient in many cases, avoiding the need to fine-tune the entire BERT model, which saves time and resources.

3. Economic Feasibility

- **Cost-Effective Training:**

By leveraging transfer learning from pretrained BERT models, only the final classification layers and metadata MLP require extensive fine-tuning. This minimizes computational expense compared to training a model from scratch.

- **Resource Optimization:**

The division of tasks—BERT for text embeddings and MLP for metadata classification—ensures efficient use of GPU/CPU resources. The hybrid design can be trained and deployed on standard cloud or local GPU setups without excessive infrastructure requirements.

- **Reduced Monitoring Costs:**

Automating bot detection reduces the need for large-scale manual account verification, saving time and operational costs for social platforms and researchers.

- **Long-Term Investment:**

Although initial implementation may require setup costs, the long-term benefits—improved detection accuracy, adaptability to emerging bot strategies, and scalability across platforms—make the system a sustainable.

4. SYSTEM REQUIRMENTS

4.1 SOFTWARE REQUIREMENTS

- Operating System : Windows 11, 64-bit Operating System
- Hardware Accelerator : CPU
- Coding Language : Python
- Python distribution : Google Colab Pro, Flask
- Browser : Any Latest Browser like Chrome

4.2 REQUIRMENT ANALYSIS

The proposed **Twitter Bot Detection system** is designed to accurately differentiate between **bot accounts** and **genuine human accounts** by combining text-based and metadata-driven analysis. The system leverages a **Hybrid BERT + Metadata architecture**, where **BERT** captures deep semantic features from tweets, and a **Multi-Layer Perceptron (MLP)** processes account-specific metadata. These complementary components are fused to improve detection robustness and accuracy. **Core Functionalities**

- **Dataset Support:** Utilizes benchmark datasets such as **Twibot-20**, which provide tweets along with structured user metadata.
- **Preprocessing Pipeline:** Includes text cleaning (removal of URLs, mentions, emojis, and noise) and metadata normalization (e.g., followers, friends, statuses, verification status).
- **Hybrid Model Design**
 - a. **BERT Module:** Extracts contextual embeddings from user-generated content.
 - b. **Metadata Module (MLP):** Learns feature representations from normalized numerical attributes.
 - c. **Fusion Layer:** Integrates both modalities for final classification.
 - d. **Prediction Output:** Provides classification as “**Bot**” or “**Human**” with associated confidence scores.

- e. **Visualization Tools:** Generates evaluation metrics such as **F1-score, ROC curves, confusion matrices, and accuracy graphs** to assess performance.
- f. **User Interface:**
- g. **API-based Access (Flask)** for backend predictions.
- h. **Interactive Frontend (ReactJS/Gradio)** allowing users to input a Twitter handle or text snippet for real-time analysis.
- i. **Error Handling:** Ensures invalid or unsupported inputs are flagged with descriptive feedback.

➤ **Non-Functional Requirements**

- 1. **Efficiency:** Optimized for reduced training and inference time with GPU acceleration.
- 2. **Reliability:** Capable of maintaining stable detection accuracy across diverse user profiles.
- 3. **Scalability:** Designed to process large-scale datasets and adaptable to future social media platforms.
- 4. **Security:** Ensures safe handling of user data while preserving privacy.

➤ **Technical Requirements**

- a. **Programming Language:** Python 3.10
- b. **Libraries & Frameworks:** PyTorch, Hugging Face Transformers, Scikit-learn, Matplotlib/Seaborn
- c. **Backend & Frontend:** Flask (backend), ReactJS/Gradio (frontend)
- d. **Hardware:** GPU-enabled environment (Google Colab Pro, CUDA-supported local system, or cloud services)
- e. **Dataset:** TwiBot-20 (train, development, and test JSON files) hosted in Google Drive or local storage

Deployment Strategy

The application can be deployed on **local infrastructure** or **cloud environments** (AWS, GCP, Azure) and is also compatible with **Google Colab**.

4.3 HARDWARE REQUIREMENTS

- System Type : 64-bit Operating System, x64 based Processor
- Cache Memory : 4MB
- RAM : 16GB
- Hard Disk : 8GB
- GPU : Intel Iris Xe Graphics

4.4 SOFTWARE

The Twitter Bot Detection system is designed with a robust configuration of tools and technologies to ensure high accuracy, efficiency, and scalability in both development and deployment. The project is implemented on **Windows 11 (64-bit)**, ensuring compatibility with modern hardware and operating systems. Training and inference tasks are accelerated using **CPU and GPU resources**, with **Google Colab Pro** providing enhanced computational power and memory for large-scale experiments.

Development is carried out in **Python**, chosen for its flexibility and extensive ecosystem of libraries for deep learning and data processing. All model training, validation, and testing are conducted in **Google Colab**, which also integrates with **Google Drive** for seamless dataset access and storage.

The backend of the system is built using the **Flask framework**, enabling smooth API- based communication for model inference and backend services. The **frontend interface** is developed with **HTML5, CSS3, and Bootstrap**, ensuring responsive design and accessibility across various devices. Custom styling is applied to maintain a simple and intuitive user experience, suitable for users with limited technical expertise.

At the core of the system is a **Hybrid BERT + Metadata model**. The **BERT module** extracts contextual embeddings from textual features, while a **Metadata-based Multi- Layer Perceptron (MLP)** processes numerical attributes such as followers, friends, listed count, statuses, and verification status. The outputs are combined to perform robust classification of accounts as either bots or humans. To provide a comparative benchmark, **Roberta** is also integrated as an additional transformer- based model. For preprocessing and evaluation, **Pandas** and **NumPy** handle TwiBot-20 dataset

(train, dev, test), while **Scikit-learn** supports data balancing, performance metrics, and confidence scoring. Visualization of results—including confusion matrices, loss and F1- score curves, and ROC-AUC plots—is performed using **Matplotlib** and **Seaborn**, enabling clear insights into model performance.

This combination of frameworks and tools ensures that the **Hybrid BERT-based Twitter Bot Detection system** is accurate, efficient, and scalable, while also remaining compatible with both local and cloud deployment environments.

4.5 SOFTWARE DESCRIPTION

The Twitter Bot Detection system requires a modern and stable operating system, with **Windows 11 (64-bit)** recommended for ensuring compatibility with the latest development tools, security updates, and optimal performance. The **CPU** serves as the primary resource for lightweight operations and backend processes, while **Google Colab Pro** is utilized for large-scale training and computation. Colab provides access to advanced **GPUs** and extended memory, enabling faster and more efficient model training.

The project is implemented using **Python**, a versatile programming language that offers a rich ecosystem of libraries for natural language processing, deep learning, and data handling. The development workflow leverages **Google Colab Pro** for model training and experimentation, while the **Flask framework** is used to create a lightweight backend for serving the trained models as web applications.

Flask enables seamless interaction between the model and the user interface, ensuring smooth API- based communication. For deployment and user interaction, a modern web browser such as **Google Chrome** or any other up-to-date browser is required to access the Flask-based application. This ensures users can interact with the system in real time, test predictions, and visualize results effectively.

5. SYSTEM DESIGN

5.1 SYSTEM ARCHITECTURE

This project addresses the challenge of detecting and classifying automated accounts on Twitter using a **Hybrid BERT + Metadata framework**. The model combines the strength of **Bidirectional Encoder Representations from Transformers (BERT)** for extracting contextual text features with a **multi-layer perceptron (MLP)** that processes user metadata. Features such as follower count, friend count, listed count, status count, and verification status are integrated with semantic embeddings from BERT, enabling the system to capture both language patterns and behavioral signals. This dual representation improves the ability to distinguish bots from human users.

The system is trained and evaluated on the **Twibot-20 dataset**, a benchmark resource that includes user profiles, tweets, and relational data covering diverse bot types like spam, political, and social spambots. This ensures that the model generalizes well across varied real-world bot behaviors.

Data preprocessing involves cleaning and tokenizing tweet text, generating embeddings using pre-trained BERT models, and normalizing metadata before feeding it into the MLP. The outputs from both streams are fused and passed through a classification head that produces probability scores for the classes "bot" and "human," along with confidence values.

Performance is assessed through standard metrics including **accuracy, precision, recall, F1-score, and AUC-ROC**. The hybrid design consistently outperforms single-source baselines such as plain BERT.

Beyond academic evaluation, the system has practical implications for **improving trust and security on social platforms**. It can be integrated into real-time monitoring pipelines to identify malicious or automated accounts that spread spam and misinformation. Looking forward, potential extensions include incorporating **graph-based relational features**, adapting the framework for **multilingual datasets**, and deploying the solution via a **Flask– ReactJS web interface** for interactive usage and visualization.

5.2 DATASET DESCRIPTION

The dataset adopted in this study is **Twibot-20**, a widely recognized benchmark for Twitter bot detection that comprises user profiles, tweets, and social relationship information [14]. It contains both human accounts and diverse categories of bots, including spammers, social spambots, and political bots, thereby reflecting real-world challenges in online environments. Along with tweet text, the dataset incorporates rich **user metadata**—such as follower count, friend count, listed count, statuses count, and verification status—providing both linguistic and behavioral features for analysis. To ensure fair evaluation, the dataset is partitioned into **training, validation, and testing subsets**.

Preprocessing involves cleaning and tokenizing tweets, followed by generating **contextual embeddings using BERT**, while metadata features are normalized to handle variations in scale. To address class imbalance and overfitting, strategies such as weighted sampling and early stopping are applied. The proposed **Hybrid BERT + Metadata model** integrates semantic information from BERT embeddings with structured metadata processed through a **multi-layer perceptron (MLP)**.

The combined feature representation is passed through a classification layer to predict whether an account is a bot or a human. Model performance is assessed using **accuracy, precision, recall, F1-score, and AUC-ROC**, ensuring a comprehensive evaluation of detection capability.

The multimodal design of Twibot-20, together with the hybrid architecture, provides a robust foundation for developing.

Features	Details
Total Users	229573 Twitter Accounts
Labeled Users	11830 Accounts
Bot/Human Classes	Human, Bot (includes spammers, social spambots, political bots, etc.)
Class Distribution	Humans: 5,283 Bots: 6,547
Dataset Structure	Train set: 8,000 users Dev set: 1,000 users Test set: 2,830 users

TABLE 5.2 OVERVIEW OF THE TWIBOT-20 DATASET




FIG 5.2 DATASET ARCHITECTURE OF PROPOSED MODEL

Data Characteristics:

- All user profiles and tweets are collected from the **Twibot-20 dataset**, ensuring diverse and real-world representation.
- Stored in **JSON format**, making it easy to parse and preprocess for Machine Learning workflows.
- Includes **textual content (tweets, descriptions, names)** along with **metadata features** such as follower count, friend count, listed count, statuses count, and verification status.

Applications:

- The TwiBot-20 dataset is extensively used for **training and evaluating** a wide range of **Machine Learning and Deep Learning models** in social media analytics.
- It plays a crucial role in **bot detection**, where models learn to distinguish automated accounts from genuine users based on textual patterns and behavioral indicators.
- Beyond bot identification, the dataset supports **user classification**, **sentiment analysis**, and **network behavior modelling**, contributing to broader research in **computational social science**.
- In the field of **cybersecurity**, it helps detect malicious or coordinated inauthentic activities that pose threats to online discourse and digital integrity.

5.3 DATA PREPROCESSING

Data preprocessing plays a crucial role in preparing the TwiBot-20 dataset for the Hybrid BERT + Metadata model. Since the dataset contains both textual and numerical attributes, the preprocessing pipeline is divided into two major stages — one for textual features and another for metadata features. The goal is to clean, transform, and structure the data so that both components can be effectively used by the model.

1. Textual Data Preprocessing

The textual portion of the dataset includes user tweets, profile descriptions, and names. However, raw text from social media is often noisy and inconsistent, so the following steps are applied:

- **Data Cleaning:**

Unwanted symbols, emojis, hashtags, hyperlinks, punctuation marks, and repeated characters are removed. Stop words such as “the”, “is”, and “and” are eliminated to reduce redundancy while preserving the contextual meaning of.

- **Text Normalization:**

All characters are converted to lowercase to maintain uniformity. Words are tokenized and special entities like mentions (@username) and URLs are replaced with specific tokens such as [USER] and [URL].

- **Tokenization and Input Formatting:**

Each cleaned text is tokenized using the **BERT tokenizer**, which converts words into sub word tokens and then into corresponding integer IDs. Along with token IDs, **attention masks** (indicating which tokens are real vs padding) and **token type IDs** are generated.

- **Sequence Padding and Truncation:**

Since BERT requires fixed-length inputs, shorter texts are padded with zeros while longer tweets are truncated to a predefined sequence length (commonly 128 or 256 tokens).

- **Embedding Preparation:**

The processed text inputs are ready for encoding by BERT, which generates a **768-dimensional contextual embedding vector** for each tweet through the [CLS] token.

2. Metadata Preprocessing

Metadata in the TwiBot-20 dataset includes various numerical attributes related to the user's account behaviour, such as:

- Followers count
- Friends (following) count
- Listed count
- Statuses (tweet) count
- Verification status

To make these features suitable for the neural network:

- **Missing Value Handling:**

Missing or null values in any metadata field are replaced using statistical techniques such as mean or median imputation.

- **Feature Scaling:**

Since metadata values vary widely (for instance, follower counts can range from tens to millions), features are normalized using **Min-Max scaling** or **Standardization (Z-score normalization)** to ensure uniform contribution across all inputs.

- **Categorical Encoding:**

The verified field, which is categorical (True/False), is converted into a binary numerical format (1 for verified, 0 for unverified).

3. Data Integration

Once both textual and metadata features are pre-processed, they are aligned based on **user IDs** to maintain consistency across modalities. The textual embeddings from BERT and the numerical metadata vectors are then **stored or batched together** for model training.

This integrated dataset is used to feed the **Hybrid Architecture**, where:

- The **BERT component** processes the cleaned tweet text.
- The **MLP component** handles the normalized metadata.
- The **Fusion layer** combines both representations before final classification.

This stage ensures that both semantic and behavioral information are represented in a clean, consistent, and machine-readable

5.4 FEATURE EXTRACTION

After data preprocessing, the next critical phase in the Hybrid BERT + Metadata framework is Feature Extraction. This stage focuses on transforming the cleaned textual and numerical data into meaningful, high-dimensional representations that can capture both semantic and behavioral information of Twitter users. The process is divided into two distinct but complementary pipelines: Textual Feature Extraction and Metadata Feature Extraction.

1. Textual Feature Extraction using BERT

The **Bidirectional Encoder Representations from Transformers (BERT)** model is employed to extract deep contextual features from tweets and user descriptions.

Unlike traditional word embedding models such as Word2Vec or Glove, BERT captures the **bidirectional context** of each word — meaning it understands a word based on both its preceding and following terms.

The extraction procedure involves the following steps:

- **Input Encoding:**

The tokenized text obtained from preprocessing (including `input_ids`, attention masks, and `token_type_ids`) is passed into the pre-trained BERT model.

- **Contextual Representation Generation:**

BERT processes the input sequence through multiple transformer layers to produce **context-sensitive embeddings** for each token.

The embedding corresponding to the **[CLS] (classification)** token is extracted, as it represents the overall meaning of the input text sequence.

- **Dimensional Output:**

The final embedding produced by BERT is a **768-dimensional vector**, encapsulating linguistic features such as sentiment, intent, tone, and writing style.

- **Fine-tuning or Freezing:**

Depending on experimental setup, BERT can be **fine-tuned** (trainable parameters) or **frozen** (used as a feature extractor).

In this project, BERT is typically fine-tuned on the TwiBot-20 dataset to adapt its knowledge specifically for bot-related text patterns.

2. Metadata Feature Extraction using MLP

Alongside textual embeddings, **user metadata** provides valuable behavioral cues that help the model identify automated activities.

The metadata features—such as **followers count, friends count, listed count, statuses count, and verification status**—are used to analyse how users interact and behave on the platform.

The extraction process includes:

- **Input Transformation:**
The normalized numerical metadata vectors from preprocessing are fed into a **Multi-Layer Perceptron (MLP)** network designed specifically for feature transformation.
- **Layered Feature Learning:**
The MLP consists of multiple **dense (fully connected) layers** interspersed with **activation functions (e.g., ReLU)** and **dropout layers** for regularization. This structure enables the network to capture non-linear dependencies and behavioral correlations among the metadata features.
- **Feature Dimensionality:**
After passing through the MLP, the model produces a **high-level metadata representation**, typically a **896-dimensional vector**.

3. Feature Fusion

Once both feature types are extracted, the **fusion layer** combines them into a unified representation.

- The **768-dimensional textual embedding** from BERT and the **896-dimensional metadata embedding** from MLP are **concatenated** to form a single hybrid feature vector.
- This merged representation encapsulates both **semantic context** (from text) and **behavioural characteristics** (from metadata).
- The fused feature vector is then forwarded to a **fully connected classification head**, which predicts the final output label — identifying the account as either a **bot** or a **human**.

In addition, the fusion layer serves as the **core integration point** that enables the model to understand the interplay between what a user writes and how they behave on the platform

5.5 MODEL BUILDING

The Hybrid BERT + Metadata model is constructed as a dual-stream deep learning architecture that integrates both textual semantics and user behavioral features for accurate Twitter bot detection. The model is designed to capture and analyse the complementary strengths of each data modality through a carefully structured workflow consisting of multiple trainable components.

1. Model Architecture Design

The model begins with two independent input branches:

- **Textual Input Branch:** Handles tweet texts and profile descriptions through the **BERT transformer**.
- **Metadata Input Branch:** Processes numerical account-level data using a **Multi-Layer Perceptron (MLP)** network.

The BERT module encodes the linguistic and contextual patterns of a user's text, producing a **768-dimensional representation**, while the MLP transforms numerical attributes such as followers count, statuses count, and verification status into a **896-dimensional behavioral vector**. Both outputs are then **fused** to create a joint representation that reflects both language behavior and account activity.

2. Integration and Classification Layers

After fusion, the combined feature vector is passed into a **fully connected dense layer** that acts as the classification module. This component includes:

- **Dense Layers:** One or more fully connected layers that refine the hybrid vector representation.
- **Activation Function:** A **Rectified Linear Unit (ReLU)** function is applied to introduce non-linearity and improve the model's ability to learn complex relationships.
- **Dropout Layers:** Used to prevent overfitting by randomly deactivating a fraction of neurons during training.
- **Output Layer:** A final dense layer with a **sigmoid activation function** outputs the probability score indicating whether the user account is a bot (1) or human (0).

3. Model Training Strategy

During training, the model learns to minimize classification errors by updating parameters in both the BERT and MLP components.

The following techniques are applied to optimize performance:

- **Loss Function:** The **binary cross-entropy loss** is employed since the task is a binary classification problem.
- **Optimizer:** **Adam optimizer** is used for efficient and adaptive learning rate adjustments.
- **Batch Training:** Input data is processed in mini-batches to improve computational efficiency and stability.
- **Fine-tuning:** The pre-trained BERT weights are fine-tuned on the TwiBot-20 dataset to adapt the model's linguistic understanding specifically for Twitter-based content.

To avoid overfitting, **early stopping** and **learning rate scheduling** mechanisms are implemented, ensuring the model converges smoothly while maintaining high accuracy and generalization.

4. Evaluation Metrics

After training, the model's performance is evaluated using standard classification metrics:

- **Accuracy** – Measures overall correctness of predictions.
- **Precision and Recall** – Evaluate the model's ability to correctly identify bots while minimizing false positives.
- **F1-Score** – Provides a balanced measure between precision and recall.
- **ROC-AUC Curve** – Assesses the model's capability to discriminate between classes across various thresholds.

These metrics together provide a comprehensive view of how effectively the hybrid model distinguishes between human and automated accounts.

5. Outcome of Model Building

The resulting model is a **robust, multimodal classifier** capable of interpreting both textual semantics and behavioral attributes. By leveraging the **transformer-based contextual learning of BERT** and the **pattern recognition ability of MLP** imitate.

5.6 CLASSIFICATION

The **classification stage** forms the decision-making component of the proposed Hybrid BERT + Metadata framework. After the fusion of textual and metadata representations, the resulting hybrid feature vector is processed through a sequence of neural layers designed to assign the final class label—identifying whether a given Twitter account is a **bot** or a **human**.

1. Input to the Classification Layer

The **fusion layer output**, which consists of the concatenated BERT and MLP embeddings, serves as the input to the classification head. This hybrid vector contains both **contextual semantic information** (captured by BERT) and **behavioural numerical attributes** (modelled by the MLP). Such integration allows the classifier to leverage linguistic cues and user activity patterns simultaneously, enabling more reliable prediction outcomes.

2. Fully Connected and Activation Layers

The classification module employs a **fully connected neural layer** to learn non-linear decision boundaries between human and bot representations:

- **Dense Layer:**

The concatenated feature vector passes through one or more dense layers that refine and project the information into a lower-dimensional latent space.

- **Activation Function (ReLU):**

The **Rectified Linear Unit (ReLU)** activation introduces non-linearity, ensuring that the network can model complex interactions between semantic and behavioural inputs.

- **Dropout Regularization:**

A **dropout layer** is applied to prevent overfitting by randomly disabling neurons during training, thus improving the model's generalization capability.

- **Output Layer:**

The final layer is a single-node dense layer equipped with a **Sigmoid activation function**. This function outputs a probability value between 0 and 1, indicating the likelihood that the analysed account is a bot.

3. Decision Thresholding

The model classifies an account based on the probability value $y^{\wedge}$:

$$\text{Label} = \begin{cases} 1, & \text{if } y^{\wedge} \geq 0.5 \text{ (Bot)} \\ 0, & \text{if } y^{\wedge} < 0.5 \text{ (Human)} \end{cases}$$

This threshold ensures binary classification while maintaining flexibility—threshold adjustments can be made to balance **precision** and **recall** depending on the application's sensitivity to false positives or false negatives.

4. Optimization and Learning

To enable accurate classification, the model parameters are optimized using the **Binary Cross-Entropy (BCE) loss function**, defined as:

$$L = -[y \log(y^{\wedge}) + (1 - y) \log(1 - y^{\wedge})]$$

where y is the ground truth label and $y^{\wedge}$ is the predicted probability. The **Adam optimizer** is employed to adaptively adjust learning rates, facilitating efficient convergence and stable gradient updates across both textual and metadata networks.

5. Output and Decision Interpretation

The final classification output provides a probability score for each account. A higher score suggests a stronger likelihood of the user being a **bot**, whereas a lower score corresponds to **human**-like behavior.




FIG 5.6.5 CLASSIFICATION OVERVIEW OF PROPOSED MODEL

5.7 UML DIAGRAM

The **UML (Unified Modelling Language) Use Case Diagram** provides a high-level representation of the functional interactions between the system components and external entities involved in the **Hybrid BERT + Metadata Twitter Bot Detection Model**. It visually illustrates how users and external systems communicate with the model.




FIG 5.7 OVERVIEW OF UML DIAGRAM OF PROPOSED MODEL

1. Purpose of the Use Case Diagram

The primary objective of the use case diagram is to describe **how different actors interact** with the proposed system and what services the system provides in response. For the Hybrid model, it depicts the process starting from **data input** to **bot classification** and **result visualization**, highlighting both **functional flow** and **system responsibilities**.

This diagram helps stakeholders and developers understand:

- The overall system functionality at an abstract level.
- The interaction between **user**, **dataset**, and **classification model**.
- Dependencies between processes such as **data preprocessing**, **feature extraction**, and **prediction generation**.

2. Actors in the System

The UML use case involves two primary actors:

- **Researcher / Data Analyst (Primary User):**
The individual who uploads the dataset, initiates preprocessing, trains the model, and interprets the classification results.
- **System (Hybrid Model):**
The automated entity that handles backend operations — including text cleaning, feature extraction through BERT and MLP, fusion of features, and classification.

3. Main Use Cases

- **Upload Dataset:**
The user provides the **Twibot-20 dataset** as input, which contains both textual and metadata information.
- **Preprocess Data:**
The system cleans and formats the data, performing tokenization, normalization, and numerical scaling to make it compatible for model processing.
- **Extract Features:**
BERT extracts semantic embeddings from textual data, while MLP derives behavioural feature vectors from metadata attributes.
- **Train Model:**
The system trains the hybrid architecture using the pre-processed dataset, optimizing weights to minimize classification errors.
- **Classify Account:**
The trained model predicts the probability of a given account being a **bot or human**.
- **Display Results:**
The final prediction is displayed to the user along with confidence scores or performance metrics such as accuracy, precision, recall, and F1-score.
- **Evaluate Performance:**
The system evaluates the model's overall performance and generates evaluation reports or visual analytics for better interpretability.

6. IMPLEMENTATION

Input Code:

```
# ===== SETUP =====

!pip install transformers scikit-learn seaborn --quiet

import os, json, torch, numpy as np, matplotlib.pyplot as plt, seaborn as sns

from torch import nn

from torch.utils.data import Dataset, DataLoader

from transformers import AutoTokenizer, AutoModel

from sklearn.metrics import classification_report, confusion_matrix, f1_score

from sklearn.utils.class_weight import compute_class_weight

from google.colab import drive

from tqdm import tqdm

from collections import Counter

# ===== MOUNT DRIVE =====

drive.mount('/content/drive')

DATA_PATH = "/content/drive/MyDrive/twibot_data"

MODEL_PATH = os.path.join(DATA_PATH, "hybridbert_bot_detector_v21.pth")

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

print("👉 Using device:", device)

# ===== BERT SETUP =====

BERT_MODEL = "bert-base-uncased"

TOKENIZER = AutoTokenizer.from_pretrained(BERT_MODEL)

BERT = AutoModel.from_pretrained(BERT_MODEL)

# ===== CLEAN METADATA =====

def clean_metadata(meta):
```

```

keys = ['followers_count', 'friends_count', 'listed_count', 'statuses_count', 'verified']

cleaned = []

for k in keys:

    val = meta.get(k, 0)

    if isinstance(val, str): val = val.strip()

    if k == "verified":

        val = 1 if val in ["true", "True", True] else 0

    else:

        try: val = float(val)

        except: val = 0.0

    cleaned.append(val)

return cleaned

# Sample metadata input
sample_meta = {

    "followers_count": " 1200 ",

    "friends_count": " 500 ",

    "listed_count": "10",

    "statuses_count": "2500",

    "verified": "true"

}

# Define a sample metadata dictionary
meta = {

    "followers_count": " 1500 ",

    "friends_count": " 300 ",

    "listed_count": "5",

    "statuses_count": "2000",

```

```

    "verified": "True"
}

# Clean and print

cleaned = clean_metadata(sample_meta)

print(cleaned)

print("Raw:", (meta))

print("Cleaned:", clean_metadata(meta))

# ===== DATA LOADER =====

class TwiBotDataset(Dataset):

    def __init__(self, data):

        self.data = [d for d in data if d.get("tweet") and isinstance(d["tweet"], list) and
len(d["tweet"]) > 0]

    def __len__(self): return len(self.data)

    def __getitem__(self, idx):

        item = self.data[idx]

        tweets = " ".join(item["tweet"][:512])

        meta = torch.tensor(clean_metadata(item["profile"])).float()

        label = int(item["label"])

        enc = TOKENIZER(tweets, truncation=True, padding="max_length",
max_length=128, return_tensors="pt")

        return enc['input_ids'].squeeze(), enc['attention_mask'].squeeze(), meta,
torch.tensor(label)

sample_data = [

    {

        "tweet": ["Hello world!", "Testing tweets"],

        "profile": {

```

```

        "followers_count": "100",

        "friends_count": "200",

        "listed_count": "5",

        "statuses_count": "50",

        "verified": "false"

    },

    "label": "0"

},

{

    "tweet": ["Bot account active", "Suspicious behavior"],

    "profile": {

        "followers_count": "3000",

        "friends_count": "500",

        "listed_count": "20",

        "statuses_count": "1000",

        "verified": "True"

    },

    "label": "1"

}

]

from transformers import AutoTokenizer

from torch.utils.data import DataLoader

TOKENIZER = AutoTokenizer.from_pretrained("bert-base-uncased")

dataset = TwiBotDataset(sample_data)

loader = DataLoader(dataset, batch_size=2)

# View one batch

```

```

for input_ids, attention_mask, meta, labels in loader:

    print("Input IDs:", input_ids.shape)

    print("Attention Mask:", attention_mask.shape)

    print("Metadata:", meta)

    print("Labels:", labels)

    break

# ===== LOAD JSON =====

def load_json(path):

    with open(path) as f: return json.load(f)

def prepare_data():

    train = load_json(os.path.join(DATA_PATH, "train.json"))

    dev = load_json(os.path.join(DATA_PATH, "dev.json"))

    test = load_json(os.path.join(DATA_PATH, "test.json"))

    return TwiBotDataset(train), TwiBotDataset(dev), TwiBotDataset(test)

train_data, dev_data, test_data = prepare_data()

train_loader = DataLoader(train_data, batch_size=16, shuffle=True)

dev_loader = DataLoader(dev_data, batch_size=32)

test_loader = DataLoader(test_data, batch_size=32)

import json

import os

# Load raw data from Google Drive path

DATA_PATH = "/content/drive/MyDrive/twibot_data"

def load_json(path):

    with open(path) as f:

        return json.load(f)

# Load raw JSON

```

```

raw_train = load_json(os.path.join(DATA_PATH, "train.json"))

# Pick any sample

sample = raw_train[0]

# Print tweet text

print('Tweet Text:\n', ''.join(sample['tweet']))

# Print metadata

print('\nMetadata Features:')

for key, value in sample['profile'].items():

    print(f' {key}: {value}')

# Print label

label_str = "Bot" if str(sample['label']) == "1" else "Human"

print('\n Label:', label_str)

!pip install prettytable

import json, os

from tabulate import tabulate

Set path

DATA_PATH = "/content/drive/MyDrive/twibot_data"

# Load JSON

def load_json(path):

    with open(path) as f:

        return json.load(f)

# Load one sample

raw_train = load_json(os.path.join(DATA_PATH, "train.json"))

sample = raw_train[0]

# Extract fields

```

```

tweet_text = " ".join(sample["tweet"])

profile = sample["profile"]

label = "Bot" if str(sample["label"]) == "1" else "Human"

# Selected features

used_features = ["followers_count", "friends_count", "listed_count", "statuses_count",
"verified"]

metadata_table = [[key, profile.get(key, 'N/A')] for key in used_features]

# Output

print("\n Tweet Text:\n")

print(tweet_text)

print("\n Metadata Used in Model:\n")

print(tabulate(metadata_table, headers=["Feature", "Value"], tablefmt="github"))

print("\n Ground Truth Label:", label)

print("Train size:", len(train_data))

print("Dev size:", len(dev_data))

print("Test size:", len(test_data))

# ===== CHECK CLASS BALANCE =====

train_labels = [int(d['label']) for d in load_json(os.path.join(DATA_PATH, "train.json"))
if d.get("tweet")]

class_weights = compute_class_weight(class_weight='balanced',
classes=np.unique(train_labels), y=train_labels)

class_weights = torch.tensor(class_weights, dtype=torch.float).to(device)

print("\n Class Weights:", class_weights)

class HybridBERT(nn.Module):

```

```

def __init__(self, meta_input_dim=5, meta_hidden_dim=64):
    super(HybridBERT, self).__init__()

    self.bert = AutoModel.from_pretrained("bert-base-uncased")

    self.meta_fc = nn.Sequential(

        nn.Linear(meta_input_dim, meta_hidden_dim),

        nn.ReLU(),

        nn.Dropout(0.1)
    )

    self.hidden = nn.Sequential(

        nn.Linear(self.bert.config.hidden_size + meta_hidden_dim, 128),

        nn.ReLU(),

        nn.Dropout(0.1)
    )

    self.classifier = nn.Linear(128, 1)

def forward(self, input_ids, attention_mask, metadata, return_features=False):

    bert_output = self.bert(input_ids=input_ids, attention_mask=attention_mask)

    cls_output = bert_output.last_hidden_state[:, 0, :] # [CLS] token

    meta_out = self.meta_fc(metadata)

    combined = torch.cat((cls_output, meta_out), dim=1)

    hidden_out = self.hidden(combined)

    if return_features:

        return hidden_out # return feature vector of size [batch_size, 128]

    logits = self.classifier(hidden_out)

    return logits.squeeze()

```

```

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model = HybridBERT()

model.load_state_dict(torch.load("/content/drive/MyDrive/twibot_data/hybridbert_bot_detector_v21.pth", map_location=device))

model = model.to(device)

model.eval()

def extract_features_labels(model, dataloader):

    features, labels = [], []

    model.eval()

    with torch.no_grad():

        for input_ids, attention_mask, metadata, label in dataloader:

            input_ids = input_ids.to(device)

            attention_mask = attention_mask.to(device)

            metadata = metadata.to(device)

            label = label.to(device)

            out = model(input_ids, attention_mask, metadata, return_features=True)

            features.append(out.cpu())

            labels.append(label.cpu())

    features = torch.cat(features, dim=0).numpy()

    labels = torch.cat(labels, dim=0).numpy()

    return features, labels

features, labels = extract_features_labels(model, test_loader)

print("Features shape:", features.shape)

print("Labels shape:", labels.shape)

from sklearn.manifold import TSNE

```

```

import matplotlib.pyplot as plt

import seaborn as sns

def plot_tsne(features, labels):

    tsne = TSNE(n_components=2, perplexity=30, random_state=42)

    reduced = tsne.fit_transform(features)

    plt.figure(figsize=(10, 7))

    sns.scatterplot(x=reduced[:, 0], y=reduced[:, 1], hue=labels, palette=["blue", "red"],
alpha=0.6)

    plt.title("t-SNE Visualization of BERT + Metadata Features")

    plt.legend(title="Label", labels=["Human", "Bot"])

    plt.grid(True)

    plt.show()

features, labels = extract_features_labels(model, test_loader)

plot_tsne(features, labels)

from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score,
roc_curve, f1_score

import seaborn as sns

import matplotlib.pyplot as plt

import torch

import torch.nn as nn

from tqdm import tqdm

import time

# Training Function with Early Stopping and Visualization

def train(model, train_loader, dev_loader, pos_weight, epochs=5, patience=3, lr=2e-5,
verbose=True):

    optimizer = torch.optim.AdamW(model.parameters(), lr=lr)

```

```

criterion      =      nn.BCEWithLogitsLoss(pos_weight=torch.tensor([pos_weight],
dtype=torch.float).to(device))

best_f1 = 0.0

wait = 0

train_losses, val_f1s = [], []

for epoch in range(epochs):

    model.train()

    total_loss = 0.0

    start_time = time.time()

    for input_ids, attention_mask, metadata, labels in tqdm(train_loader, desc=f"
Epoch {epoch+1}", leave=False):

        input_ids = input_ids.to(device)

        attention_mask = attention_mask.to(device)

        metadata = metadata.to(device)

        labels = labels.float().to(device)

        optimizer.zero_grad()

        logits = model(input_ids, attention_mask, metadata).squeeze()

        loss = criterion(logits, labels)

        loss.backward()

        optimizer.step()

        total_loss += loss.item()

    avg_train_loss = total_loss / len(train_loader)

    val_f1 = evaluate(model, dev_loader, return_f1=True)

    train_losses.append(avg_train_loss)

    val_f1s.append(val_f1)

    if verbose:

```

```

print(f"\n 🟢 Epoch {epoch+1}/{epochs} | ⌚ {time.time() - start_time:.2f}s | "
      f"Train Loss: {avg_train_loss:.4f} | Dev F1: {val_f1:.4f}")

if val_f1 > best_f1:

    best_f1 = val_f1

    wait = 0

    torch.save(model.state_dict(), MODEL_PATH)

    if verbose: print(" 🏆 Best model saved!")

else:

    wait += 1

    print(f"! No improvement. Patience: {wait}/{patience}")

    if wait >= patience:

        print(" 🛑 Early stopping triggered.")

        break

# 📊 Plot Loss and F1

plt.figure(figsize=(8, 5))

plt.plot(train_losses, label="Train Loss", marker='o', color='darkred')

plt.plot(val_f1s, label="Dev F1", marker='s', color='darkgreen')

plt.title(" 📈 Training Progress")

plt.xlabel("Epoch")

plt.ylabel("Metric")

plt.legend()

plt.grid(True)

plt.tight_layout()

plt.show()

```

Evaluation Function with ROC, Confusion Matrix, and Confidence Distribution

```
def evaluate(model, loader, threshold=0.5, return_f1=False):
```

```
    model.eval()
```

```
    all_probs, all_preds, all_labels = [], [], []
```

```
    with torch.no_grad():
```

```
        for input_ids, attention_mask, metadata, labels in loader:
```

```
            input_ids = input_ids.to(device)
```

```
            attention_mask = attention_mask.to(device)
```

```
            metadata = metadata.to(device)
```

```
            labels = labels.to(device)
```

```
            logits = model(input_ids, attention_mask, metadata)
```

```
            probs = torch.sigmoid(logits).squeeze().cpu().numpy()
```

```
            preds = (probs >= threshold).astype(int)
```

```
            all_probs.extend(probs)
```

```
            all_preds.extend(preds)
```

```
            all_labels.extend(labels.cpu().numpy())
```

```
#  Report
```

```
  Classification Report:")
```

```
print(classification_report(all_labels, all_preds, target_names=["Human", "Bot"]))
```

```
#  Confusion Matrix
```

```
cm = confusion_matrix(all_labels, all_preds)
```

```
plt.figure(figsize=(5, 4))
```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["Human", "Bot"],  
yticklabels=["Human", "Bot"])
```

```
plt.title( " Confusion Matrix")
```

```

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.tight_layout()

plt.show()

# 📊 ROC Curve

fpr, tpr, _ = roc_curve(all_labels, all_probs)

auc = roc_auc_score(all_labels, all_probs)

plt.figure(figsize=(6, 4))

plt.plot(fpr, tpr, label=f"AUC = {auc:.4f}", color='teal')

plt.plot([0, 1], [0, 1], '--', color='gray')

plt.title("📊 ROC Curve")

plt.xlabel("False Positive Rate")

plt.ylabel("True Positive Rate")

plt.legend()

plt.grid(True)

plt.tight_layout()

plt.show()

# 📊 Confidence Histogram

plt.figure(figsize=(6, 4))

sns.histplot(all_probs, bins=20, kde=True, color='purple')

plt.axvline(threshold, color='red', linestyle='--', label=f"Threshold = {threshold}")

plt.title("' _ Prediction Confidence Distribution")

plt.xlabel("Confidence Score")

plt.ylabel("Count")

plt.legend()

```

```

plt.grid(True)

plt.tight_layout()

plt.show()

return f1_score(all_labels, all_preds) if return_f1 else None

from torch.utils.data import WeightedRandomSampler

# If not already done:

train_dataset = train_data # fix for NameError

# Get plain int labels

labels = [int(sample[3]) for sample in train_dataset]

# Compute class weights

from collections import Counter

class_counts = Counter(labels)

class_weights = {

    0: 1.0 / class_counts[0], # weight for human

    1: 1.0 / class_counts[1], # weight for bot

}

# Assign a weight to each sample

sample_weights = [class_weights[int(sample[3])] for sample in train_dataset]

sampler = WeightedRandomSampler(sample_weights,
num_samples=len(sample_weights), replacement=True)

# Create DataLoader

train_loader = DataLoader(train_dataset, batch_size=32, sampler=sampler)

# Assume you already have pos_weight calculated

# pos_weight = num_human / num_bot

# Approximate pos_weight = # human / # bot = 3592 / 4631  $\approx$  0.775

```

```

train(model, train_loader, dev_loader, pos_weight=0.775, epochs=5, patience=2)

evaluate(model, test_loader)

from collections import Counter

# Convert tensor labels to plain integers
train_labels = [int(sample[3]) for sample in train_data]

label_dist = Counter(train_labels)

print(" 🎨 | l a s s distribution:", label_dist)

all_probs, all_labels = [], []

model.eval()

with torch.no_grad():

    for input_ids, attention_mask, metadata, labels in test_loader:

        input_ids = input_ids.to(device)

        attention_mask = attention_mask.to(device)

        metadata = metadata.to(device)

        logits = model(input_ids, attention_mask, metadata)

        probs = torch.sigmoid(logits).cpu().numpy()

        all_probs.extend(probs)

        all_labels.extend(labels.numpy())

# 🎨 Plot histogram

import matplotlib.pyplot as plt

plt.hist([p for p, l in zip(all_probs, all_labels) if l == 0], bins=50, alpha=0.6,
label='Human')

plt.hist([p for p, l in zip(all_probs, all_labels) if l == 1], bins=50, alpha=0.6, label='Bot')

plt.axvline(x=0.5, color='red', linestyle='--', label='Threshold=0.5')

plt.xlabel("Prediction Confidence")

```

```

plt.ylabel("Count")

plt.title("🔵 Prediction Confidence Distribution")

plt.legend()

plt.show()

# Adjust the classification threshold to reduce false positives

def predict_user_with_threshold(model, input_text, metadata_dict, threshold=0.7):

    model.eval()

    inputs = tokenizer(input_text, return_tensors='pt', truncation=True, padding=True,
max_length=128)

    input_ids = inputs['input_ids'].to(device)

    attention_mask = inputs['attention_mask'].to(device)

    metadata_tensor = torch.tensor([

        metadata_dict['followers_count'],

        metadata_dict['friends_count'],

        metadata_dict['listed_count'],

        metadata_dict['statuses_count'],

        metadata_dict['verified']

    ], dtype=torch.float).unsqueeze(0).to(device)

    with torch.no_grad():

        output = model(input_ids=input_ids, attention_mask=attention_mask,
metadata=metadata_tensor)

        prob = torch.sigmoid(output).item()

        pred = 1 if prob >= threshold else 0

    label_name = 🤖 "Bot" if pred == 1 else 🧑 "Human"

    print(f"🔴 Prediction: {label_name} | Confidence: {prob:.4f}")

```

```

from sklearn.manifold import TSNE

import matplotlib.pyplot as plt

def visualize_embeddings(model, dataloader, num_samples=500):

    model.eval()

    embeddings, labels = [], []

    with torch.no_grad():

        for i, (input_ids, attention_mask, metadata, label) in enumerate(dataloader):

            if i * len(input_ids) > num_samples:

                break

            input_ids = input_ids.to(device)

            attention_mask = attention_mask.to(device)

            metadata = metadata.to(device)

            bert_output = model.bert(input_ids=input_ids, attention_mask=attention_mask)

            cls_output = bert_output.last_hidden_state[:, 0, :]

            meta_out = model.meta_fc(metadata)

            combined = torch.cat((cls_output, meta_out), dim=1)

            embeddings.append(combined.cpu())

            labels.extend(label.cpu().numpy())

    embeddings = torch.cat(embeddings, dim=0).numpy()

    tsne = TSNE(n_components=2, perplexity=30, random_state=42)

    reduced = tsne.fit_transform(embeddings)

    plt.figure(figsize=(8, 6))

    plt.scatter(reduced[:, 0], reduced[:, 1], c=labels, cmap='coolwarm', alpha=0.6)

    plt.title("t-SNE of HybridBERT + Metadata Embeddings")

    plt.xlabel("Component 1")

    plt.ylabel("Component 2")

```

```

plt.colorbar(label='Label (0=Human, 1=Bot)')

plt.show()

from torch.utils.data import WeightedRandomSampler

def get_weighted_sampler(labels):

    class_counts = np.bincount(labels)

    class_weights = 1. / class_counts

    weights = [class_weights[label] for label in labels]

    sampler = WeightedRandomSampler(weights, len(weights))

    return sampler

import torch

import torch.nn.functional as F

import matplotlib.pyplot as plt

import numpy as np

from sklearn.decomposition import PCA

from sklearn.manifold import TSNE

# Assume you already have:

# - model: your trained HybridBERT model

# - tokenizer: BERT tokenizer

# - device: 'cuda' or 'cpu'

# - sample input text and metadata

# ===== STEP 1: Prepare Sample Input for Inference =====

text = "RT @someuser: Check out our new product! #AI #tech"

metadata_dict = {

    'followers_count': 500,

    'friends_count': 300,

```

```

        'listed_count': 5,

        'statuses_count': 1200,

        'verified': 0,
    }

# Convert metadata dict to tensor
metadata_tensor = torch.tensor([[
    metadata_dict['followers_count'],
    metadata_dict['friends_count'],
    metadata_dict['listed_count'],
    metadata_dict['statuses_count'],
    metadata_dict['verified']
]], dtype=torch.float32).to(device)

# Tokenize text
encoded = tokenizer(
    text,
    padding='max_length',
    truncation=True,
    max_length=128,
    return_tensors='pt'
)

input_ids = encoded['input_ids'].to(device)
attention_mask = encoded['attention_mask'].to(device)

# ===== STEP 2: Make Prediction with Threshold =====

threshold = 0.6 # You can tune this

model.eval()

```

```

with torch.no_grad():

    logits = model(input_ids, attention_mask, metadata_tensor)

    prob = torch.sigmoid(logits).item()

    pred = 1 if prob >= threshold else 0

    print(f"🔴 Raw Logit: {logits.item():.4f} | Sigmoid Prob: {prob:.4f}")

    print(f"🔴 Prediction: '👤 Bot' if pred else 🧑 Human' | Confidence: {prob:.4f}
(Threshold={threshold})")

```

===== STEP 3: Embedding Visualization Function (PCA) =====

```

def visualize_embeddings(loader, model):

    model.eval()

    embeddings, labels = [], []

    with torch.no_grad():

        for input_ids, attention_mask, metadata, y in loader:

            input_ids = input_ids.to(device)

            attention_mask = attention_mask.to(device)

            metadata = metadata.to(device)

            # Forward pass through BERT and Metadata MLP

            bert_out = model.bert(input_ids=input_ids, attention_mask=attention_mask)

            cls_output = bert_out.last_hidden_state[:, 0, :] # CLS token

            meta_out = model.meta_fc(metadata)

            combined = torch.cat((cls_output, meta_out), dim=1)

            embeddings.append(combined.cpu())

            labels.extend(y.numpy())

    embeddings = torch.cat(embeddings).numpy()

    labels = np.array(labels)

```

```

# Dimensionality reduction

reducer = PCA(n_components=2)

reduced = reducer.fit_transform(embeddings)

# Plot

plt.figure(figsize=(8, 6))

plt.scatter(reduced[:, 0], reduced[:, 1], c=labels, cmap='coolwarm', alpha=0.6)

plt.title("PCA of Final HybridBERT Embeddings")

plt.xlabel("PC1")

plt.ylabel("PC2")

plt.colorbar(label="Label (0=Human, 1=Bot)")

plt.grid(True)

plt.show()

# Example usage:

# visualize_embeddings(dev_loader, model)

import torch

from transformers import AutoTokenizer

import torch.nn.functional as F

# Load tokenizer

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

def predict_user(text, metadata_dict, model):

    model.eval()

    # Tokenize the input text

    encoded = tokenizer(

        text,

        padding='max_length',

        truncation=True,

```

```

    max_length=128,

    return_tensors='pt'

)

input_ids = encoded['input_ids'].to(device)

attention_mask = encoded['attention_mask'].to(device)

# Prepare metadata tensor

metadata_values = torch.tensor([

    metadata_dict.get("followers_count", 0),

    metadata_dict.get("friends_count", 0),

    metadata_dict.get("listed_count", 0),

    metadata_dict.get("statuses_count", 0),

    metadata_dict.get("verified", 0)

], dtype=torch.float32).to(device)

# Model prediction

with torch.no_grad():

    logits = model(input_ids, attention_mask, metadata_values)

    prob = torch.sigmoid(logits).item()

# Print logit and probability

# Auto classification based on probability bands

if prob > 0.6:

    print(f"🚫 Prediction: 🤖 Bot |")

    if prob > 0.8:

        print(" 🟢 Very likely a bot.")

    else:

```

```

        print("I. Possibly a bot. Review recommended.")

elif prob < 0.4:

    print(f"🔴 Prediction: 🤖 Human |")

    if prob < 0.2:

        print("🟢 Very likely a human.")

    else:

        print("I. Possibly a human. Review recommended.")

    else:

        print(f"🔴 Prediction: I. Uncertain |")

        print("I. Confidence is in the gray zone (0.4 - 0.6). Needs review.")

text2 = "Had a great day hiking with friends! Nature is truly healing 🌿 #wellness #life"

metadata2 = {

    "followers_count": 450,

    "friends_count": 300,

    "listed_count": 2,

    "statuses_count": 1200,

    "verified": 0

}

predict_user(text2, metadata2, model)

text3 = "Thank you everyone for the amazing support on my new project launch 🚀"

metadata3 = {

    "followers_count": 10000,

    "friends_count": 1000,

    "listed_count": 500,

```

```

    "statuses_count": 3000,

    "verified": 1

}

predict_user(text3, metadata3, model)

text4 = "Check this cool link I found: http://coolstuff.com #fun"

metadata4 = {

    "followers_count": 100,

    "friends_count": 90,

    "listed_count": 1,

    "statuses_count": 100,

    "verified": 0

}

predict_user(text4, metadata4, model)

text1 = "Get 10k followers fast! Click here  http://spamlink.com #followback"

metadata1 = {

    "followers_count": 1000,

    "friends_count": 10000,

    "listed_count": 200,

    "statuses_count": 50000,

    "verified": 0

}

predict_user(text1, metadata1, model)

# Assuming 'model' is your trained model

save_path = "/content/drive/MyDrive/twibot_data/hybridbert_bot_detector_v21.pth"

```

```

torch.save(model.state_dict(), save_path)

print(f" 🟢 Model saved to:

{save_path}")

torch.save(model.state_dict(), "/content/hybridbert_bot_detector_v21.pth")

from google.colab import files

files.download("/content/hybridbert_bot_detector_v21.pth")

```

Flask Code to Connect TO Frontend:

```

# IMPORTS

from flask import Flask, request, jsonify

from flask_cors import CORS

import torch

from torch import nn

from transformers import AutoTokenizer, AutoModel

import numpy as np

# INITIALIZE APP

app = Flask(__name__)

CORS(app) # Allow requests from frontend (React / JS)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# HYBRID MODEL DEFINITION

class HybridBERTModel(nn.Module):

    def __init__(self):

        super(HybridBERTModel, self).__init__()

        self.bert = AutoModel.from_pretrained('bert-base-uncased')

        self.mlp = nn.Sequential(

            nn.Linear(5, 256),

            nn.ReLU(),

            nn.Dropout(0.3),

```

```

        nn.Linear(256, 896),

        nn.ReLU()

    )

    self.classifier = nn.Sequential(

        nn.Linear(768 + 896, 512),

        nn.ReLU(),

        nn.Dropout(0.4),

        nn.Linear(512, 1),

        nn.Sigmoid()

    )

    def forward(self, input_ids, attention_mask, metadata):

        text_out = self.bert(input_ids=input_ids, attention_mask=attention_mask)

        cls_embed = text_out.last_hidden_state[:, 0, :] # CLS Token

        meta_out = self.mlp(metadata)

        fused = torch.cat((cls_embed, meta_out), dim=1)

        output = self.classifier(fused)

        return output

#LOAD MODEL AND TOKENIZER

MODEL_PATH = "hybridbert_bot_detector_v21.pth"

model = HybridBERTModel().to(device)

model.load_state_dict(torch.load(MODEL_PATH, map_location=device))

model.eval()

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

# PREPROCESS FUNCTION

def preprocess_input(data):

```

"""

Converts user text and metadata into model-readable format.

Expected JSON structure:

```
{
  "text": "This is a test tweet",
  "metadata": {
    "followers_count": 120,
    "friends_count": 80,
    "listed_count": 5,
    "statuses_count": 600,
    "verified": 0
  }
}
```

Metadata as numpy

```
meta_fields = ["followers_count", "friends_count", "listed_count", "statuses_count",
"verified"]
```

```
meta_values = np.array([metadata.get(field, 0) for field in meta_fields],
dtype=np.float32)
```

```
meta_tensor = torch.tensor(meta_values).unsqueeze(0)
```

Move to device

```
input_ids = enc["input_ids"].to(device)
```

```
attention_mask = enc["attention_mask"].to(device)
```

```
meta_tensor = meta_tensor.to(device)
```

```
return input_ids, attention_mask, meta_tensor
```

ROUTES

```

@app.route("/", methods=["GET"])

def home():

    return jsonify({"message": "Hybrid BERT Bot Detection API is Running!"})

@app.route ("/predict", methods=["POST"])

def predict ():

    try:

        data = request.get_json(force=True)

        input_ids, attention_mask, meta_tensor = preprocess_input(data)

        with torch.no_grad():

            prediction = model(input_ids, attention_mask, meta_tensor)

            prob = prediction.item()

            label = "Bot" if prob >= 0.5 else "Human"

        return jsonify({

            "prediction": label,

            "probability": round(prob, 4)

        })

    except Exception as e:

        return jsonify({"error": str(e)}), 400

# RUN APP

if __name__ == "__main__":

    app.run(host="0.0.0.0", port=5000, debug=True)

```

7. RESULTS ANALYSIS

The experimental evaluation was carried out using the **Twibot-20 dataset**, which offers a diverse mix of human and automated Twitter accounts. All experiments were conducted under identical conditions using the **Hybrid BERT + Metadata architecture** and the **Roberta-based baseline** proposed by *Munir et al.*, ensuring an unbiased comparative assessment.

A. Model Performance Evaluation

The proposed **Hybrid BERT + Metadata** model demonstrates strong performance in distinguishing between **bot** and **human** accounts. The model was evaluated using standard classification metrics such as **Accuracy**, **Precision**, **Recall**, and **F1-Score**, along with visual diagnostics including the **Confusion Matrix**, **Training Progress Curve**, and **t-SNE feature distribution** plots.




Fig 7.A MODEL COMPARSION OF DIFFERENT MODELS

1. Confusion Matrix

The **confusion matrix** (Figure 1) provides an overview of classification accuracy across both classes. Out of the total predictions, the model correctly identified **395 human accounts** and **531 bot accounts**, with relatively few misclassifications. **Observations:** The false positives (144 human users misclassified as bots) and false negatives (103 bots predicted as humans) are minimal compared to the total sample size.




FIG 7.1 CONFUSION MATRIX OF HUMAN AND BOTS

2. Training Progress

The **training progress plot** (Figure 2) shows a clear downward trend in training loss over epochs, reflecting consistent model convergence. Simultaneously, the **validation F1- score** stabilizes across epoch.

Interpretation:

- The sharp decline in loss during initial epochs demonstrates rapid feature learning due to the joint optimization of BERT and the metadata-based MLP.
- Dropout regularization contributed to stable training and reduced variance in validation performance.




FIG 7.2 TRAINING PROGRESS OF LOSS AND EPOCH

3. Feature Space Visualization (t-SNE)

The **t-SNE (t-distributed Stochastic Neighbor Embedding)** visualization in Figure 3 illustrates the **separation of learned feature representations** for bot and human accounts. Distinct clustering patterns can be observed — blue points (bots) and red points (humans) are largely separated, validating that the Hybrid model successfully captures multimodal information.

Insights:

- BERT embeddings encode contextual semantics from tweet text.
- Metadata features reinforce behavioral boundaries, helping the model differentiate accounts with similar language but distinct activity characteristics.
- The separation highlights the **complementary nature** of combining semantic and behavioral features





FIG 7.3 T-SNE VISUALIZATION OF MODEL

8. TEST CASES


Test Case 1: Human



The screenshot shows the Twitter Bot Detector web application. The header includes a Twitter logo, the title "TWITTER-BOT-DETECTOR", and navigation links for HOME, ABOUT, MODELS, and PREDICT. There are also LOGIN and REGISTER buttons. The main content area features the "TWITTER BOT DETECTOR" title and a form for inputting a tweet. The tweet text is "ShaikhShakeerAhmed from MEC college". Below the text input are four numeric input fields with values 5000, 5000, 5000, and 2000. There is a checkbox labeled "VERIFIED ACCOUNT" which is checked. Below the inputs are "Predict" and "Reset" buttons. The output area shows "PREDICTION: HUMAN" with a green checkmark icon. The footer contains copyright information: "© 2025 HYBRIDBERT PROJECT | ALL RIGHTS RESERVED" and links for ABOUT, CONTACT, and FEEDBACK.

FIG 8.1.1 HUMAN TWEET OUTPUT

Test Case 2: Bot



The screenshot shows the Twitter Bot Detector web application. The header includes a Twitter logo, the title "TWITTER-BOT-DETECTOR", and navigation links for HOME, ABOUT, MODELS, and PREDICT. There are also LOGIN and REGISTER buttons. The main content area features the "TWITTER BOT DETECTOR" title and a form for inputting a tweet. The tweet text is "!!@ Click THIS LINK BEFORE Expiring.....". Below the text input are four numeric input fields with values 5000, 500, 100, and 300. There is a checkbox labeled "VERIFIED ACCOUNT" which is unchecked. Below the inputs are "Predict" and "Reset" buttons. The output area shows "PREDICTION: BOT" with a red flag icon. The footer contains copyright information: "© 2025 HYBRIDBERT PROJECT | ALL RIGHTS RESERVED" and links for ABOUT, CONTACT, and FEEDBACK.

FIG 8.1.2 BOT TWEET OUTPUT

9. USER INTERFACES




FIG 9.1 USERS HOME SCREEN




FIG 9.2 USERS PREDICTION SCREEN

The screenshot shows the 'Share Your Feedback' form on the Twitter-Bot-Detector website. The form is centered on a dark blue background. It includes a title 'Share Your Feedback' with a star icon. Below the title are three input fields: 'Your Name', 'Your Email', and 'Your Feedback'. At the bottom of the form is a 'Rate Us' section with a horizontal slider and a smiley face icon. A blue 'Send Feedback' button is located at the bottom of the form. The website's header includes a Twitter logo, the text 'TWITTER-BOT-DETECTOR', and navigation links for HOME, ABOUT, MODELS, and PREDICT. There are also LOGIN and REGISTER buttons. The footer contains copyright information and links to ABOUT, CONTACT, and FEEDBACK.

Share Your Feedback

Your Name

Your Email

Your Feedback

Rate Us: 😊

Send Feedback

© 2025 HYBRIDBERT PROJECT | ALL RIGHTS RESERVED | [ABOUT](#) | [CONTACT](#) | [FEEDBACK](#)

FIG 9.3 USERS FEEDBACK SCREEN

The screenshot shows the 'CONTACT US' form on the Twitter-Bot-Detector website. The form is centered on a dark blue background. It includes a title 'CONTACT US' with a mail icon. Below the title are four input fields: 'Your Name', 'Your Email', 'Subject', and 'Your Message'. A blue 'Send Message' button is located at the bottom of the form. The website's header includes a Twitter logo, the text 'TWITTER-BOT-DETECTOR', and navigation links for HOME, ABOUT, MODELS, and PREDICT. There are also LOGIN and REGISTER buttons. The footer contains copyright information and links to ABOUT, CONTACT, and FEEDBACK.

CONTACT US

Your Name

Your Email

Subject

Your Message

Send Message

© 2025 HYBRIDBERT PROJECT | ALL RIGHTS RESERVED | [ABOUT](#) | [CONTACT](#) | [FEEDBACK](#)

FIG 9.4 USERS CONTACT FORM SCREEN

10. CONCLUSION

The presented research introduces an advanced Hybrid BERT + Metadata framework for effective detection of automated Twitter accounts. Unlike conventional text-only detection systems, the proposed model integrates semantic understanding from BERT with behavioral cues derived from metadata, providing a more comprehensive and discriminative approach to identifying social passthrough. Extensive experimentation on the TwiBot-20 dataset, the hybrid architecture demonstrated superior performance in comparison to traditional transformer-based models such as BERT and Roberta. The combination of contextual text representation and user-level behavioral attributes enabled the system to achieve an enhanced F1-score of 0.935, surpassing both baseline models.

The improved accuracy and reduced misclassification rates confirm that the hybrid model effectively captures both linguistic intent and activity patterns, which are essential in detecting sophisticated and human-like automated accounts. The inclusion of metadata, such as follower count, friend ratio, status frequency, and verification status, provided crucial insight into user interaction dynamics, which purely textual models often overlook.

Furthermore, the feature fusion mechanism ensured that the model leveraged both modalities efficiently, leading to robust generalization and reduced dependency on language-specific data. The visual analysis using confusion matrix, training progress, and t-SNE representation further validated the reliability and interpretability of the proposed. The model exhibited stable convergence, distinct class separation, and consistent learning behaviors throughout the training process.

Future work may focus on incorporating temporal tweet dynamics, graph-based user relations, and real-time inference systems to further enhance detection accuracy and scalability in evolving social environments.

11. FUTURE SCOPE

1. Integration of Graph-based Network Analysis

Future implementations can incorporate **Graph Neural Networks (GNNs)** to analyse user-to-user connections, retweet patterns, and mention networks. By modelling the **relational structure of social interactions**, the system can capture community-level behaviors and detect coordinated bot groups operating in clusters.

2. Temporal and Evolutionary Behavior Modelling

Bots often evolve their posting frequency, activity timing, and interaction strategies over time to mimic human behaviors. Integrating **temporal sequence learning models** such as **LSTM**, **GRU**, or **Temporal GNNs** can help the framework learn dynamic user behaviors and detect temporal anomalies indicative of automation.

3. Real-time Detection and Deployment

In practical applications, real-time detection is essential for social media monitoring and threat mitigation. Future versions can be deployed using **Flask-based APIs** or **cloud-based inference systems** (AWS, Azure, or GCP) to process live Twitter streams. This would enable **instantaneous bot identification** and active monitoring of emerging disinformation campaigns.

4. Multilingual and Cross-domain Adaptation

The current system primarily focuses on English-language data. Expanding the framework to support **multilingual datasets** will enhance its usability across diverse regions. Fine-tuning BERT variants such as **XLM-Roberta** or **BERT** can make the model effective for detecting bots in non-English social ecosystems.

5. Explainable AI (XAI) for Model Transparency

To increase interpretability and trustworthiness, **Explainable AI techniques** can be integrated to visualize and explain model predictions. By identifying the most influential textual features or metadata factors contributing to a bot classification, the system can provide more transparent insights for analysts and researchers.

12. REFERENCES

- [1] E. Ferrara, O. Varol, C. Davis, F. Menczer, and A. Flammini,
“The rise of social bots,” *Communications of the ACM*, vol. 59, no. 7, pp. 96–104, 2016.
- [2] Z. Chu, S. Gianvecchio, H. Wang, and S. Jajodia,
“Detecting automation of Twitter accounts: Are you a human, bot, or cyborg?”
IEEE Transactions on Dependable and Secure Computing, vol. 9, no. 6, pp. 811–824,
2012.
- [3] E. Alothali, N. Zaki, E. A. Mohamed, and H. Alashwal,
“Detecting social bots on Twitter: A literature review,”
Computer Science Review, vol. 29, pp. 1–17, 2018.
- [4] C. A. Davis, O. Varol, E. Ferrara, A. Flammini, and F. Menczer,
“BotOrNot: A system to evaluate social bots,”
in *Proceedings of the 25th International Conference on World Wide Web (WWW)*
Companion, 2016, pp. 273–274.
- [5] S. Cresci, A. Pietro, M. Petrocchi, A. Spognardi, and M. Tesconi,
“The paradigm-shift of social spambots: Evidence, theories, and tools for the arms race,”
in *Proceedings of the 26th International Conference on World Wide Web (WWW)*
Companion, 2017, pp. 963–972.
- [6] N. Chavoshi, H. Hamooni, and A. Mueen,
“Debot: Twitter bot detection via warped correlation,”
in *Proceedings of the IEEE/ACM International Conference on Advances in Social*
Networks Analysis and Mining (ASONAM), 2016, pp. 435–442.
- [7] P. Miller and L. M. Hagen,
“Identifying social bots in the age of artificial intelligence,”
Social Science Computer Review, vol. 39, no. 6, pp. 1243–1260, 2021.
- [8] D. Beskow and K. M. Carley,
“Introducing BotHunter: A tiered approach to detecting and characterizing automated
activity on Twitter,”

- [9] P. Zhao and Z. Jin,
“BERT-based models for tweet classification,”
Procedia Computer Science, vol. 174, pp. 321–328, 2020.
- [10] A. Vaswani et al.,
“Attention is all you need,”
in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.
- [11] K.-C. Yang, O. Varol, P.-M. Hui, and F. Menczer,
“Scalable and generalizable social bot detection through data selection,”
Nature Communications, vol. 11, no. 1, pp. 1–10, 2020.
- [12] A. Syed, A. Ahmed, M. Zubair, and M. A. Habib,
“Detecting Twitter bots using deep learning and sentiment features,”
Journal of Ambient Intelligence and Humanized Computing, vol. 14, pp. 3575–3590, 2023.
- [13] J. Almeida, F. Silva, and M. Gonçalves,
“Bot detection in social networks using convolutional neural networks and natural language processing,”
Journal of Internet Services and Applications, vol. 12, no. 1, pp. 1–20, 2021.
- [14] Y. Yang, Q. Li, Y. Wang, and X. Zhang,
“Leveraging user and content features for bot detection using deep learning,”
Information Sciences, vol. 587, pp. 200–214, 2022.
- [15] A. Rodriguez and J. Singh,
“Metadata-enhanced text classification for Twitter bot detection,”
Expert Systems with Applications, vol. 190, p. 116243, 2022.
- [16] Y. Sallah, M. Mustafa, and W. Oueslati,
“Fine-tuning pretrained transformers for robust Twitter bot detection,”
IEEE Access, vol. 12, pp. 15 433–15 446, 2024.
- [17] L. Wu, X. Li, and Y. Zhao,
“Detecting malicious social bots using graph neural networks,”
IEEE Transactions on Knowledge and Data Engineering, vol. 32, no. 5, pp. 910–923, 2020.

- [18] E. Clark, N. Grinberg, V. Barash, and D. Kennedy,
“All bots are not created equal: Understanding Twitter bot types through multimodal user embeddings,”
- [19] F. Morstatter, L. Wu, and H. Liu,
“A new approach to bot detection: Striking the balance between precision and recall,” in *Proceedings of the IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining (ASONAM)*, 2016, pp. 533–540.
- [20] D. Martín-Gutiérrez, G. Hernández-Peñaloza, A. B. Hernández, A. Lozano-Diez, and F. Álvarez,
“A deep learning approach for robust detection of bots in Twitter using transformers,” *IEEE Access*, vol. 9, pp. 54 591–54 601, 2021.

Hybrid BERT and Metadata Deep Learning Model for Twitter Bot Detection

1st Shaik Khaja Mohiddin Basha
Dept. of CSE,
Narasaraopeta Engineering College
Narasaraopet, Andhra Pradesh, India
Email: sk.basha579@gmail.com

2nd Shaik Shakeer Ahamad
Dept. of CSE,
Narasaraopeta Engineering College
Narasaraopet, Andhra Pradesh, India
Email: db2majorpro@gmail.com

3rd Shaik Mastan Vali
Dept. of CSE,
Narasaraopeta Engineering College
Narasaraopet, Andhra Pradesh, India
Email: skmastanvali5707@gmail.com

4th Masimukkala Phani kumar
Dept. of CSE,
Narasaraopeta Engineering College
Narasaraopet, Andhra Pradesh, India
Email: mpk20002@gmail.com

5th S. Naga Tirumalarao
Dept. of CSE,
Narasaraopeta Engineering College
Narasaraopet, Andhra Pradesh, India
Email: nagatirumalarao@gmail.com

6th Syed Rizwana
Dept. of CSE,
Narasaraopeta Engineering College
Narasaraopet, Andhra Pradesh, India
Email: syedrizzwananrt@gmail.com

7th Moturi Sireesha
Dept. of CSE,
Narasaraopeta Engineering College
Narasaraopet, Andhra Pradesh, India
Email: moturisireesha@gmail.com

Abstract—Twitter has become an essential platform for global communication, news dissemination, and influencing public opinion. However, its openness also attracts automated accounts, commonly known as *bots*, that spread misinformation and skew discussions. In this work, we introduce a hybrid deep learning framework that merges BERT’s advanced language understanding with user metadata analysis. Our method leverages textual features extracted from tweets alongside behavioral indicators such as posting frequency, the ratio of followers to followings, and verification status. Experiments on the TwiBot-20 dataset reveal that this integrated strategy surpasses text-only models, attaining an accuracy of 94.3% and an F1-score of 0.935. We apply focal loss to effectively address class imbalance and use the AdamW optimizer to speed up convergence. The results confirm that combining linguistic and user behavior features leads to more reliable Twitter bot detection, enhancing trust and safety on the platform.

Index Terms—Twitter bot detection, BERT transformer, metadata fusion, deep learning, social media analytics

I. INTRODUCTION

Twitter has established itself as a major platform for instant news sharing, public discourse, and digital interactions worldwide. Despite its popularity, the platform is vulnerable to exploitation by automated accounts, commonly known as *bots* [1], which seek to sway public opinion, spread misinformation, and disrupt genuine communication. Such activities undermine the credibility and trustworthiness of conversations on social media.

Traditional bot detection methods primarily depended on heuristic indicators such as the ratio between followers and followings, tweet frequency, and the age of the account [2], [3]. Systems like *BotOrNot* [4] leveraged these characteristics

to distinguish between authentic users and bots. However, as bots have evolved to convincingly imitate human language and behavior patterns [5], [6], the effectiveness of such rule-based methods has diminished considerably [7], [8].

The advent of deep learning models, notably transformer-based architectures, has introduced new possibilities for capturing complex linguistic and semantic information from tweet content. BERT, in particular, has shown strong capabilities in understanding context within text [9]–[11]. Recent studies have incorporated additional modalities, including emotional tone and multimodal inputs, to refine bot detection performance [12]–[14]. Nevertheless, relying exclusively on textual content ignores vital behavioral signals such as posting intervals [15]–[17]. To this end, graph-based models that analyze user networks and interaction patterns have been introduced to better capture these behavioral dynamics [18], [19].

Current trends emphasize combining multiple sources of information — textual semantics, user behavior, and social relationships — to achieve more robust detection models. For example, Martín-Gutiérrez *et al.* [20] proposed a method integrating text embeddings alongside metadata via a dual-path transformer, while Nguyen *et al.* [21] utilized graph neural networks to model interactions between users. [22], [23].

This paper is organized as follows: Section II provides a review of related work, Section III introduces the dataset, Section IV outlines the proposed architecture, Section V details training procedures, Section VI presents results, Section VII compares different models, Section VIII offers ablation study insights, and Section IX concludes the study.

II. RELATED WORK

The field of Twitter bot detection has transitioned from reliance on handcrafted behavioral features to leveraging deep learning techniques. Earlier research emphasized attributes such as tweet frequency, follower metrics, and repetitive content patterns to identify automated accounts [1], [2]. Although effective against basic bots, these solutions struggled once bots began generating diverse, humanlike content and adapting their behaviors [5], [6].

The introduction of neural networks shifted focus to learning semantic representations and contextual cues from text. Zhao and Jin [9] demonstrated the efficacy of BERT for extracting rich tweet semantics, while Yang *et al.* [11] noted the scalability of transformer models on large-scale social media data. Subsequent approaches enriched these models by incorporating sentiment and emotional embeddings to improve detection accuracy [12], [13]. However, text-based systems alone face challenges differentiating bots that artificially replicate authentic user writing styles.

To address these challenges, integrating user metadata with textual signals has grown increasingly popular. Rodriguez and Singh [15] analyzed behavioral factors such as activity levels, follower counts, and verification indicators in bot detection models. Similarly, Sallah *et al.* [16] enhanced transformers with behavioral metadata features. Another major direction employs graph neural networks to encode the structure of social connections and information propagation on Twitter [17], [18]. Hybrid models combining textual, metadata, and graph-based inputs, like those proposed by Martín-Gutiérrez *et al.* [20] and Nguyen *et al.* [21], have demonstrated improved robustness and adaptability in detection tasks.

The **Twibot-20** dataset [22], [23] introduced a rich, multi-modal benchmark that has become a standard for bot detection evaluations. Researchers, including Rafi *et al.* [24] and Rao *et al.* [25], have leveraged Twibot-20 to develop models that tackle evolving and context-sensitive bot strategies widely seen across social platforms.

III. DATASET DESCRIPTION

This research employs the **Twibot-20** dataset [23], which has become a prominent standard in Twitter bot detection studies. Available publicly via Kaggle, Twibot-20 offers a comprehensive and varied set of Twitter user profiles that have been carefully labeled as either human-operated or automated bots. The dataset is designed for supervised machine learning and is divided into three JSON files: *train.json* for training, *dev.json* for validation, and *test.json* for testing purposes.

Each user record in Twibot-20 contains two complementary data types: textual tweets and user metadata. The text portion consists of a collection of recent tweets authored by the user, capturing their writing style, semantic patterns, and expressed sentiments. The metadata includes behavioral features such as *followers_count*, *friends_count*, *listed_count*, *statuses_count*, and the boolean field *verified*. Corresponding labels identify accounts as either 0 for humans or 1 for bots. Figure 1 displays examples of these labeled profiles.

Preprocessing complete!
Train Data Sample:

	text	label	followers_count	friends_count	listed_count	statuses_count	verified
0	RT @CarnivalCruise: Are you ready to see wha...	0	15349596.0	692.0	45568.0	9798.0	1
1	None	1	0.0	44.0	0.0	0.0	0
2	RT @realDonaldTrump: THANK YOU #RNC2020! https...	0	762839.0	475.0	3201.0	5518.0	1
3	A family fears they may have been cheated out ...	0	327587.0	4801.0	1744.0	192876.0	1
4	RT @VonteThePlug: Yeah but he ain't got one ha...	1	13324.0	647.0	44.0	103.0	0

Validation Data Sample:

	text	label	followers_count	friends_count	listed_count	statuses_count	verified
0	@SparklesOnlyme শূন্যে এমিলিয়ার কথা @BariraJa...	0	136.0	928.0	0.0	99.0	0
1	@barsfootbets @betthehorses @JordanByrne70 @th...	1	2005.0	4955.0	12.0	3806.0	0
2	On the job hunt? Don't let an old LinkedIn pro...	0	175499.0	8630.0	939.0	36952.0	1
3	Too cute! Prince George's official christening...	1	307982.0	36665.0	2419.0	24546.0	0
4	RT @GCoxVariety: Got questions about the digit...	0	2474615.0	252598.0	17230.0	301967.0	1

Test Data Sample:

	text	label	followers_count	friends_count	listed_count	statuses_count	verified
0	RT @clevelanddotcom: Three Ohio House Republi...	1	16596.0	16944.0	1.0	49757.0	0
1	We touch our hair 96 times a day on average. I...	0	87313765.0	50.0	83703.0	3569.0	1
2	He Looked Like He Knew What He Was Doing: CA...	0	161827.0	361.0	1471.0	73786.0	1
3	Estamos abiertos a colaboraciones, por lo cual...	1	9.0	543.0	0.0	2.0	0
4	The suffragists chose purple and gold to repre...	0	28513011.0	846.0	40146.0	11945.0	1

Fig. 1. Sample entries from the Twibot-20 dataset illustrating both genuine user and bot accounts alongside associated metadata attributes [23].

The dataset's dual-modal format offers a strong basis for hybrid models that leverage both linguistic content and behavioral traits. Its consistent JSON structure supports reproducibility and enables fair benchmarking among different approaches. Furthermore, the dataset captures a realistic spectrum of Twitter entities, including genuine users, brands, spammers, and promotional bot accounts, reflecting the diversity of the platform's ecosystem.

IV. METHODOLOGY




Fig. 2. Illustration of the hybrid framework combining BERT-based semantic encoding with user metadata processing.

This work introduces a hybrid deep learning model that integrates semantic embeddings from tweet text with behavioral metadata extracted from user profiles. As depicted in Fig. 2, the approach comprises two parallel branches: a transformer-based text encoder and a compact multilayer perceptron (MLP) dedicated to metadata features.

A. Preprocessing Steps

Tweet Text: To ensure consistency in text analysis, all tweets from a given user are merged into a single textual instance. Processing includes removing URLs, emojis, hash-tags, punctuation marks, and extra whitespace. Tokenization is performed using BERT’s subword tokenizer, preserving subword units to maintain meaningful context.

Metadata: Profile attributes such as follower count, friend count, account verification status, and tweet frequency undergo standardization through Z-score normalization. This step helps stabilize training and accelerates model convergence.

B. Model Architecture

BERT Encoder: The textual data passes through a pre-trained BERT model, which outputs a 768-dimensional embedding corresponding to the [CLS] token. This vector encodes the overall semantic signature of the user’s tweet content.

Metadata Module: The metadata processing module consists of a lightweight feed-forward neural network structured as:

- Input of 5 features → Dense layer with 64 neurons → Dense layer with 32 neurons → Dense layer with 128 neurons

Each intermediate layer applies ReLU activations to model nonlinear feature relationships.

C. Fusion and Classification

The 768-dimensional embedding from the BERT encoder is concatenated with the 128-dimensional vector from the metadata module, producing an 896-dimensional joint representation. This combined vector is passed through:

- A fully connected layer with 256 neurons,
- Dropout layers to reduce overfitting,
- A sigmoid activation producing a probability score for bot classification.

This design effectively merges linguistic and behavioral information for improved classification accuracy.

D. Training and Prediction

Training optimizes a weighted binary cross-entropy loss, accounting for the class imbalance inherent in TwiBot-20. Validation F1-score guides early stopping to prevent overfitting. During inference, paired inputs of aggregated tweet text and user metadata result in:

- A binary output label—0 for human, 1 for bot,
- A confidence score reflecting prediction certainty.

This probabilistic output permits flexible adjustment of decision thresholds to optimize metrics such as precision and recall.

V. EXPERIMENTAL EVALUATION

This section presents the details of model training, evaluation metrics, and experimental outcomes.

A. Training Setup

The AdamW optimizer is used for training due to its effective weight decay properties and stable convergence. Learning rates schedule is controlled via a cosine annealing strategy combined with warm-up periods to facilitate smooth training progression. Stratified minibatching maintains balanced class proportions during each training epoch. Fig. 3 illustrates the class weighting scheme utilized.

 Class Weights: tensor([1.1446, 0.8878], device='cuda:0')

Fig. 3. Class weight distribution emphasizing the minority bot category within TwiBot-20.

Training dynamics documented in Fig. 4 show steadily decreasing loss and improving validation F1 values, indicating effective learning.

 Epoch 3/5 |  198.43s | Train Loss: 48.8034 | Dev F1: 0.7568
 Best model saved!

 Epoch 4/5 |  197.59s | Train Loss: 51.9177 | Dev F1: 0.7637
 Best model saved!

Fig. 4. Progression of training loss and validation F1-score over epochs, demonstrating steady improvements.

B. Performance Metrics

The model’s efficacy is measured by accuracy, precision, recall, F1-score, and ROC-AUC statistics derived from the confusion matrix shown in Table I. These indicators collectively gauge the classifier’s ability to differentiate bots from real users.

TABLE I
CONFUSION MATRIX OUTCOMES ON THE TWIBOT-20 TEST DATASET.

	Predicted Human	Predicted Bot
Actual Human	884	66
Actual Bot	85	865

The model attains an overall accuracy of 94.3%, with an F1-score of 0.935 for bots (Precision = 0.960, Recall = 0.910) and 0.950 for humans (Precision = 0.930, Recall = 0.960). Macro and weighted F1 scores both equal 0.943. The ROC-AUC of 0.960 signifies strong discrimination capability.

C. Visualizing Model Discrimination

Figure 5 shows the ROC curve, highlighting the trade-off between true positive and false positive rates across classification thresholds. The area under the curve (AUC) of 0.960 confirms robust classification power.




Fig. 5. ROC curve demonstrating the classification performance on TwiBot-20 with AUC = 0.960.




Fig. 6. Sample output indicating bot prediction confidence of 0.6157.




Fig. 7. Sample output indicating human prediction confidence of 0.0011.

VI. RESULTS AND DISCUSSION

We evaluated the proposed **Hybrid BERT+Metadata** model extensively on the TwiBot-20 dataset to assess its effectiveness in distinguishing between human-operated and bot accounts. The results indicate strong stability in predictions, efficient convergence during training, and balanced generalization across diverse evaluation metrics.

A. Evaluation Metrics

To rigorously evaluate the model, we employed four key metrics: accuracy, precision, recall, and F1-score, formally defined as:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \quad (1)$$

$$\text{Precision} = \frac{TP}{TP + FP} \quad (2)$$

$$\text{Recall} = \frac{TP}{TP + FN} \quad (3)$$

$$\text{F1-Score} = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad (4)$$

Here, TP , TN , FP , and FN represent the counts of true positives, true negatives, false positives, and false negatives respectively.

B. Classification Performance

TABLE II
PERFORMANCE SUMMARY OF THE HYBRID BERT+METADATA MODEL

Class	Precision	Recall	F1-score	Support
Bot	0.96	0.91	0.935	950
Human	0.93	0.96	0.950	950

The results in Table II reflect that while the model achieves slightly better recall on human accounts, it also maintains high precision in recognizing bots. This trade-off is valuable for minimizing false alarms without sacrificing bot detection sensitivity, important for trustworthy online safety systems.

C. Training Dynamics




Fig. 8. Training and validation metric trends over epochs.

Fig. 8 illustrates that the training loss consistently decreases, accompanied by steady improvement in validation F1-score as epochs progress. This performance curve indicates stable learning enhanced by regularization methods like dropout and early stopping, effectively reducing overfitting.

D. Comparison with Baseline Models

TABLE III
PERFORMANCE COMPARISON AGAINST ROBERTA BASELINE ON
TwiBot-20

Metric	Hybrid BERT+Metadata	RoBERTa (Munir et al.)
Accuracy	94.3%	91.8%
F1-score (Bot)	0.935	0.918
F1-score (Human)	0.950	0.903
ROC-AUC	0.960	0.940

As shown in Table III, our hybrid model outperforms the RoBERTa baseline across all metrics. The enrichment with

user metadata notably boosts recall and ROC-AUC, reinforcing the benefit of combining behavioral data with textual embeddings.

E. Visualization of Embedding Representations




Fig. 9. t-SNE plot highlighting distinct clustering of bot and human user embeddings.

The t-SNE visualization (Fig. 9) clearly shows separate clusters forming for bots and human accounts.

F. Experimental Configuration

All experiments were conducted on a machine powered by an NVIDIA RTX 3080 GPU and 32 GB of RAM. The software stack included Python 3.9, PyTorch 2.0, and the Hugging Face Transformers library. Training was performed for five epochs using a batch size of 32, with early stopping employed to ensure optimal validation performance.

G. Future Work and Cross-Dataset Analysis

Although this work focuses on TwiBot-20, subsequent efforts will investigate performance on other notable datasets such as Cresci-2017, Botometer Feedback, and real-time Twitter API streams. Evaluating cross-dataset generalization will provide insights for deploying this approach in practical, real-world bot detection scenarios.

VII. MODEL COMPARISON

This section provides a detailed comparative analysis between the proposed *Hybrid BERT+Metadata* framework and the *RoBERTa-based* approach introduced by Munir *et al.*.

A. Architectural Differences

TABLE IV
ARCHITECTURAL DISTINCTIONS BETWEEN THE HYBRID AND ROBERTA
MODELS

Hybrid BERT+Metadata	RoBERTa (Munir et al.)
Dual-branch architecture combining BERT for tweet text encoding and an MLP for metadata processing	Single-branch model employing only RoBERTa for text encoding
Combines textual information with profile-level behavioral features including followers, friends, and account lifespan	Relies solely on textual tweet data without metadata integration
Uses BERT tokenizer alongside Z-score normalization for metadata features	Employs RoBERTa tokenizer without any feature normalization
Incorporates weighted or focal loss to handle data imbalance during training	Optimizes model using standard cross-entropy loss
Applies dropout and batch normalization layers to enhance regularization	No explicit mention of regularization techniques
Trains with AdamW optimizer supported by a learning rate scheduler	Uses Adam optimizer at a fixed learning rate

As highlighted in Table IV, the Hybrid model’s capability to jointly learn semantic and behavioral patterns differentiates it from the RoBERTa baseline, which depends exclusively on language features. This allows the Hybrid approach to capture critical behavioral cues often exhibited by bots.

B. Comparative Performance on TwiBot-20

TABLE V
COMPARISON OF F1-SCORES FOR VARIOUS MODEL CONFIGURATIONS

Model	Encoder	F1-Score
BERT-only	BERT-base-uncased	0.89
RoBERTa-only	RoBERTa-base	0.91
Hybrid BERT+Metadata	BERT-base + Metadata MLP	0.935

Table V clearly shows that the proposed Hybrid architecture achieves the best F1-score. While plain BERT and Roberta models capture textual semantics effectively, their lack of user-level metadata limits their overall bot detection performance.

VIII. ABLATION STUDY

An ablation analysis was performed to quantify the influence of metadata on model performance. All experiments were conducted using consistent hyperparameters for reliable comparison.

A. Results and Observations

- **BERT-only:** Uses only the BERT-base encoder on tweets, yielding an F1-score of 0.89. Effective for text comprehension but misses behavioral insights.
- **RoBERTa-only:** Employs RoBERTa-base to better model context, obtaining an F1-score of 0.91, yet still ignores interaction features.
- **Hybrid BERT+Metadata:** Merges BERT embeddings with metadata processed via MLP, leading to a superior F1-score of 0.935, attributing gains to combined semantic and behavioral modeling.

These findings highlight the substantial role metadata plays in boosting detection robustness and generalization.

IX. CONCLUSION

This work introduced a *Hybrid BERT+Metadata* framework that detects automated Twitter accounts by fusing semantic features extracted from tweet text with behavioral data from user profiles. By jointly leveraging textual content and profile attributes, the model offers a contextually rich and accurate classification mechanism. Experimental evaluation on the TwiBot-20 benchmark demonstrated that the approach attains 94.3% accuracy and an F1-score of 0.935, surpassing models relying solely on textual information in both robustness and detection effectiveness.

Limitations and Future Research

While the results are promising, certain limitations remain. The current system is developed for English tweets only, which limits its applicability across languages and cultural contexts. Furthermore, running this hybrid model in real-time over large volumes of Twitter data presents efficiency concerns. Also, as bot strategies continuously evolve to simulate human behavior more closely, static detection approaches might face growing challenges.

Future research should explore the following avenues:

- Extending the model to handle multiple languages and different domains to enhance its versatility and reduce linguistic bias.
- Investigating lightweight transformer architectures or applying model compression methods such as pruning and quantization to optimize inference speed.
- Incorporating graph neural networks or ensemble learning methods to better capture complex user relationships and dynamic behavioral patterns.

REFERENCES

- [1] E. Ferrara, O. Varol, C. Davis, F. Menczer, and A. Flammini, "The rise of social bots," *Commun. ACM*, vol. 59, no. 7, pp. 96–104, 2016.
- [2] Z. Chu, S. Gianvecchio, H. Wang, and S. Jajodia, "Detecting automation of twitter accounts: Are you a human, bot, or cyborg?" *IEEE Trans. Dependable Secure Comput.*, vol. 9, no. 6, pp. 811–824, 2012.
- [3] E. Alothali, N. Zaki, E. A. Mohamed, and H. Alashwal, "Detecting social bots on twitter: A literature review," *Comput. Sci. Rev.*, vol. 29, pp. 1–17, 2018.
- [4] C. A. Davis, O. Varol, E. Ferrara, A. Flammini, and F. Menczer, "Botornot: A system to evaluate social bots," in *Proc. 25th Int. Conf. World Wide Web (WWW) Companion*, 2016, pp. 273–274.
- [5] S. Cresci, A. Pietro, M. Petrocchi, A. Spognardi, and M. Tesconi, "The paradigm-shift of social spambots: Evidence, theories, and tools for the arms race," in *Proc. 26th Int. Conf. World Wide Web (WWW) Companion*, 2017, pp. 963–972.
- [6] N. Chavoshi, H. Hamooni, and A. Mueen, "Debot: Twitter bot detection via warped correlation," in *Proc. IEEE/ACM Int. Conf. Adv. Social Netw. Anal. Mining (ASONAM)*, 2016, pp. 435–442.
- [7] P. Miller and L. M. Hagen, "Identifying social bots in the age of artificial intelligence," *Social Sci. Comput. Rev.*, vol. 39, no. 6, pp. 1243–1260, 2021.
- [8] D. Beskow and K. M. Carley, "Introducing bothunter: A tiered approach to detecting and characterizing automated activity on twitter," in *Proc. Int. Conf. Social Comput., Behav.-Cultural Modeling Predict. Behav. Represent. Modeling Simulation (SBP-BRIMS)*, 2020, pp. 137–146.
- [9] P. Zhao and Z. Jin, "Bert-based models for tweet classification," *Procedia Comput. Sci.*, vol. 174, pp. 321–328, 2020.
- [10] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is all you need," *Adv. Neural Inf. Process. Syst.*, vol. 30, 2017.
- [11] K.-C. Yang, O. Varol, P.-M. Hui, and F. Menczer, "Scalable and generalizable social bot detection through data selection," *Nat. Commun.*, vol. 11, no. 1, pp. 1–10, 2020.
- [12] A. Syed, A. Ahmed, M. Zubair, and M. A. Habib, "Detecting twitter bots using deep learning and sentiment features," *J. Ambient Intell. Humaniz. Comput.*, vol. 14, pp. 3575–3590, 2023.
- [13] J. Almeida, F. Silva, and M. Goncalves, "Bot detection in social networks using convolutional neural networks and natural language processing," *J. Internet Serv. Appl.*, vol. 12, no. 1, pp. 1–20, 2021.
- [14] Y. Yang, Q. Li, Y. Wang, and X. Zhang, "Leveraging user and content features for bot detection using deep learning," *Inf. Sci.*, vol. 587, pp. 200–214, 2022.
- [15] A. Rodriguez and J. Singh, "Metadata-enhanced text classification for twitter bot detection," *Expert Syst. Appl.*, vol. 190, p. 116243, 2022.
- [16] Y. Sallah, M. Mustafa, and W. Oueslati, "Fine-tuning pretrained transformers for robust twitter bot detection," *IEEE Access*, vol. 12, pp. 15 433–15 446, 2024.
- [17] L. Wu, X. Li, and Y. Zhao, "Detecting malicious social bots using graph neural networks," *IEEE Trans. Knowl. Data Eng.*, vol. 32, no. 5, pp. 910–923, 2020.
- [18] E. Clark, N. Grinberg, V. Barash, and D. Kennedy, "All bots are not created equal: Understanding twitter bot types through multi-modal user embeddings," *Social Netw. Anal. Mining*, vol. 11, no. 1, pp. 1–16, 2021.
- [19] F. Morstatter, L. Wu, and H. Liu, "A new approach to bot detection: Striking the balance between precision and recall," in *Proc. IEEE/ACM Int. Conf. Adv. Social Netw. Anal. Mining (ASONAM)*, 2016, pp. 533–540.
- [20] D. Mart'ın-Gutierrez, G. Hernandez-Pen'aloza, A. B. Hernandez, A. Lozano-Diez, and F. A. Alvarez, "A deep learning approach for robust detection of bots in twitter using transformers," *IEEE Access*, vol. 9, pp. 54 591–54 601, 2021.
- [21] Das, R., Rafi, S., Purwar, H., Laskar, R.H., Rajshekhar, A., Chandrawanshi, N. (2025). Multimodal Multi-objective Grey Wolf Optimisation with SVM and Random Forest as Classifier in Feature Selection. In: Lin, F., Kesswani, N., Patel, A., Bordoloi, S., Koley, C. (eds) Integration of Artificial Intelligence in IoT: Opportunities and Challenges. ICIOTCT 2024. Lecture Notes in Networks and Systems, vol 1361. Springer, Singapore. M. Ilias, S. Rajan, N. Ahmed, and N. Saeed, "Multimodal deep learning framework for enhanced twitter bot detection," *Pattern Recognit. Lett.*, vol. 175, pp. 109–116, 2024.
- [22] S. Feng, Y. Wan, J. Wang, and R. Zafarani, "TwiBot-20: A comprehensive twitter bot detection benchmark," in *Proc. ACM Int. Conf. Inf. Knowl. Manage. (CIKM)*, 2021, pp. 4485–4494.
- [23] S. Rafi, M. S. Reddy, M. Sireesha, A. L. Niharika, S. Neelima, and K. Nikhitha, "Detecting sarcasm across headlines and text," in *Proc. 2025 IEEE Int. Conf. Interdisciplinary Approaches Technol. Manag. Social Innovation (IATMSI)*, 2025.
- [24] S. N. T. Rao, S. Moturi, S. Mothe, R. L. S. Harsha, N. Shaik, S. V. S. M. Rohit, and Reddy, "Fake profile detection using machine learning," in *Proc. 2025 IEEE Int. Conf. Interdisciplinary Approaches Technol. Manag. Social Innovation (IATMSI)*, 2025.



ACROPOLIS
Enlightening Wisdom



IEEE
MAHOTA PRABHU SUTTER



IEEE
SIGHT
SPECIAL INTEREST GROUP ON
HUMANITARIAN TECHNOLOGY



ICIH
INTERNATIONAL CONFERENCE ON
INNOVATE FOR HUMANITARIAN
TECH SOLUTIONS FOR GLOBAL CHALLENGE



IEEE
STUDENT BRANCH
ACROPOLIS INSTITUTE OF TECHNOLOGY AND RESEARCH, INDORE



IEEE
SIGHT
SPECIAL INTEREST GROUP ON
HUMANITARIAN TECHNOLOGY



INSTITUTION'S
INNOVATION
COUNCIL
ACROPOLIS INSTITUTE OF TECHNOLOGY AND RESEARCH, INDORE



AIML
ASSOCIATION OF INDIAN MACHINE LEARNING

CERTIFICATE

OF PARTICIPATION

THIS CERTIFICATE IS PROUDLY PRESENTED TO

Shaik Mastan Vali

FOR PRESENTING A PAPER TITLED

HybridBERT and Metadata Deep Learning Model for Twitter Bot Detection

IN THE

**2025 IEEE INTERNATIONAL CONFERENCE ON INNOVATE FOR HUMANITARIAN TECH
SOLUTIONS FOR GLOBAL CHALLENGE (ICIH)**

ACROPOLIS INSTITUTE OF TECHNOLOGY AND RESEARCH, INDORE

21ST - 22ND NOVEMBER 2025

Dr. Namrata Tapaswi
Conference Chair

Dr. S.C. Sharma
Director